

An Adaptive Filling Strategy for the Multi-Processor Cup Game

William Kuszmaul¹, Alek Westover¹

¹Massachusetts Institute of Technology

Abstract

The problem of scheduling tasks on p processors so that no task ever gets too far behind is often described as a game with cups and water. In the p -processor cup game on n cups, there are two players, a filler and an emptier, that take turns adding and removing water from a set of n cups. In each turn, the filler adds p units of water to the cups, placing at most 1 unit of water in each cup, and then the emptier selects p cups to remove up to 1 unit of water from. The emptier's goal is to minimize the backlog, which is the height of the fullest cup.

The p -processor cup game has been studied in many different settings, dating back to the late 1960's. Kuszmaul established that a greedy emptier never lets backlog exceed $O(\log n)$. Kuszmaul also gave a construction that achieves backlog $\Omega(\log(n-p))$. For small p , e.g. $p < n/2$, it is clear that $\Theta(\log n)$ is the optimal backlog, but for very large p tight bounds are not known.

We give an adaptive filling strategy for the p -processor cup game on n cups that achieves backlog $\Omega(\log n)$. Combined with Kuszmaul's work, this implies that $\Theta(\log n)$ is the optimal backlog for the multi-processor cup game.

1 Introduction

The p -processor cup game on n cups is a multi-round game with two players, the **filler** and the **emptier**, that take turns adding and removing water from the cups. In particular, on each round the filler distributes p units of fill arbitrarily amongst the cups, subject only to the restriction that the filler cannot place more than 1 unit of fill into any given cup. Next, the emptier chooses p cups and removes at most 1 unit of water from each of these. The **backlog** is the height of the fullest cup; the emptier aims to minimize backlog while the filler aims to maximize backlog.

Kuszmaul showed that in the p -processor cup game on n cups the greedy emptying strategy—i.e. empty from the p fullest cups—never lets backlog exceed $O(\log n)$. Kuszmaul also provided a filling strategy that achieves backlog $\Omega(\log(n-p))$. For small p , e.g. $p \leq n/2$, the upper-bound and lower-bound match, but for large p the bounds do not match.

Many suspected that the $O(\log n)$ upper bound could be reduced to $O(\log(n-p))$ even for large p ; however, we show that this is not the case.

2 Preliminaries

Let $H_n = 1/1 + 1/2 + \dots + 1/n$ denote the n -th harmonic number. It is well known that $H_n = \Theta(\log n)$, and in fact $H_n = \ln n + \Theta(1)$.

The **mass** of the cups is the sum the fills of all of the cups.

We remark that we consider only **adaptive** fillers, i.e. fillers that can observe the state of the cups. The emptier can achieve better bounds on backlog by using randomization, in which case the emptier is not allowed to view the state of the cups, and is called **oblivious**.

3 Lower Bound on Backlog

Theorem 1. *There is a filling strategy for the p -processor cup game on n cups that achieves backlog $\Omega(\log n)$ against any emptier.*

Proof. Kuszmaul's construction shows that there is a filling strategy that achieves backlog $\Omega(\log(n-p))$. If $p \leq n - \sqrt{n}$, then $n-p \geq \sqrt{n}$, so $\log(n-p) \geq \frac{1}{2} \log(n)$, hence $\Omega(\log(n-p)) = \Omega(\log(n))$. Thus the result holds for $p \leq n - \sqrt{n}$; we proceed to consider the case where $p > n - \sqrt{n}$.

We give a filling strategy, that we call **IncMass**, with running time $O(p/(n-p))$, that increases the mass of the cups in n rounds by at least $1/2$, provided that backlog starts below $O(\log n)$. The filler's strategy for achieving backlog $\Omega(\log n)$ is to repeatedly use IncMass as many as $\Theta(n \log n)$ times, terminating if backlog $\Omega(\log n)$ is ever achieved (after which the filler changes to the strategy of placing a unit of water in the p fullest cups on each round). The filler's strategy either terminates before finishing, which only happens if it has achieved backlog $\Omega(\log n)$, or does not terminate early, in which case the mass of the cups has increased by $\Omega(n \log n)$, implying that which means that average fill and also backlog have increased by at least $\Omega(\log n)$. We now describe IncMass.

Let S denote the set of cups. The filler will maintain a set $U \subset S$ throughout the algorithm. The algorithm's procedure will ensure that once a cup enters U its fill never decreases for the rest of the process (by placing 1 unit of water in such a cup each round). Furthermore, $|U|$ will increase by $n-p$ at each iteration of the process. U is initialized to \emptyset . For each of $\lfloor (p+1)/(n-p) \rfloor$ steps the filler will:

1. Distribute $p - |U|$ water equally among the cups in $S \setminus U$ (thus, each such cup receives $\frac{p-|U|}{n-|U|}$ fill)
2. Distribute $|U|$ water equally among the cups in U (thus, each such cup receives 1 fill)

Then the emptier must chose p cups to empty from, and hence $n-p$ cups to **neglect**, i.e. not empty from. Let N be the set of neglected cups on this step, with $|N| = n-p$. The emptier adds all cups in $N \setminus U$ to U , and then adds the $(n-p) - |N \setminus U|$ fullest cups in $S \setminus U$ to U .

The number of cups in $S \setminus U$ decreases by $n-p$ on one of the $\lfloor (p+1)/(n-p) \rfloor$ steps, so in the end we have $|S \setminus U| \geq p+1$.

Claim 1. *Any cup $c \in S \setminus U$ at the end of this process has 0 fill, provided that the backlog started as at most $O(\log n)$.*

Proof. On the i -th step of the filler's process the fill of a cup in $S \setminus U$ increases by $(p - |U|)/(n - |U|)$, and then decreases by 1 on this round, resulting in a net change of

$$-\frac{n-p}{n - (n-p) \cdot i}.$$

The total amount that the fill of c has changed since the start of the filler's process is simply the sum of the amounts that the fill changes on each step, which is

$$\sum_{i=0}^{\lfloor \frac{p+1}{n-p} \rfloor - 1} \frac{n-p}{n - (n-p) \cdot i}. \quad (1)$$

We aim to show that (1) is $\Omega(\log n)$; combined with the fact that the backlog started as $O(\log n)$ this will imply that c has fill 0.

For $p = n-1$ (1) is simply the difference of harmonic numbers; this prompts the idea of lower bounding (1) by a difference of harmonic numbers. We lower bound the i -th term in the sum by:

$$\frac{n-p}{n - i(n-p)} = \sum_{j=0}^{n-p-1} \frac{1}{n - i(n-p)} \geq \sum_{j=n-i(n-p)}^{n-(i-1)(n-p)-1} \frac{1}{j}.$$

Applying this to (1) we get a difference of harmonic numbers:

$$\sum_{i=0}^{\lfloor \frac{p+1}{n-p} \rfloor - 1} \sum_{j=n-i(n-p)}^{n-(i-1)(n-p)-1} \frac{1}{j}.$$

This is now a difference of harmonic numbers. When $i = 0, j = n - (i - 1)(n - p) - 1$ we get the smallest term in the sum:

$$\frac{1}{n + (n - p - 1)}.$$

When $i = \lfloor (p + 1)/(n - p) \rfloor - 1, j = n - i(n - p)$ we get the largest term in the sum, which is at least

$$\frac{1}{n - \left(\frac{p+1}{n-p} - 1 \right) (n - p)} = \frac{1}{2(n - p) - 1}.$$

Thus, c has lost fill at least

$$H_{2n-p-1} - H_{2(n-p-1)} \geq \Omega\left(\log \frac{p}{n-p}\right).$$

Because $p > n - \sqrt{n}$ we have that c has lost fill at least

$$\Omega\left(\log \frac{n - \sqrt{n}}{\sqrt{n}}\right) \geq \Omega(\log n),$$

as desired. Because the backlog, and hence fill of c , started less than $O(\log n)$ by assumption, c must now have 0 fill. \square

There are now (at least) $p + 1$ cups with fill 0. The final step of the filler's procedure is to add $1/(p + 1)$ fill to $p + 1$ cups with fill 0. The emptier now must waste resources on emptying one of these cups with fill below 1. In total on this round the filler adds p units of fill to the cups, and then the emptier removes at most $p - 1 + 1/(p + 1) \leq p - 1/2$ water to the cups. Hence the mass increases by at least $1/2$. As mass is monotonically increasing, this increase of $\frac{1}{2}$ in mass persists, as desired.

We have shown that IncMass can increase the mass given that backlog starts below $O(\log n)$. As shown previously, by repeatedly using IncMass we can achieve backlog $O(\log n)$.

We remark that the running time of IncMass is $O(p/(n - p))$ which satisfies

$$\sqrt{n} \leq \frac{p}{n - p} \leq n.$$

The running time of the full strategy is

$$O\left(\frac{p}{n - p} n \log n\right) \leq O(n^2 \log n).$$

\square

4 Conclusion

Our analysis shows that $\Theta(\log n)$ is a tight bound on backlog in the p -processor cup game on n cups with an adaptive filler.

However, many open questions remain. The bounds on backlog in the multi-processor cup game with an oblivious filler have *not* yet been made tight. In particular, there is a lower bound of $\Omega(\log \log n)$, and an upper bound of $O(\log \log n + \log p)$; for $p \leq \log n$ the bounds are tight, but for large p they are not.

Further, analysis of the variant of the game where p can change, potentially with resource augmentation, is also of great interest.