

Variable-Processor Cup Games

William Kuszmaul^{*1} and Alek Westover^{†1}

¹MIT

kuszmaul@mit.edu, alek.westover@gmail.com

Abstract

In the *cup game* two players, the *filler* and the *emptier*, take turns adding and removing water from cups, subject to certain constraints. In the classic *p*-processor cup game the filler distributes p units of water among the n cups with at most 1 unit of water to any particular cup, and the emptier chooses p cups to remove at most one unit of water from. Analysis of the cup game is important for applications in processor scheduling, buffer management in networks, quality of service guarantees, and deamortization.

We investigate a new variant of the classic *p*-processor cup game, which we call the *variable-processor cup game*, in which the resources of the emptier and filler are variable. In particular, in the variable-processor cup game the filler is allowed to change p at the beginning of each round. Although the modification to allow variable resources seems small, we show that it drastically alters the game.

We construct an adaptive filling strategy that achieves backlog $\Omega(n^{1-\varepsilon})$ for any constant $\varepsilon > 0$ of our choice in running time $2^{O(\log^2 n)}$. This is enormous compared to the upper bound of $O(\log n)$ that holds in the classic *p*-processor cup game! We also present a simple adaptive filling strategy that is able to achieve backlog $\Omega(n)$ in extremely long games: it has running time $O(n!)$.

Furthermore, we demonstrate that this lower bound on backlog is tight: using a novel set of invariants we prove that a greedy emptier never lets backlog exceed $O(n)$.

We also construct an oblivious filling strategy that achieves backlog $\Omega(n^{1-\varepsilon})$ for $\varepsilon > 0$ constant of our choice in time $2^{O(\log^2 n)}$ against any “greedy-like” emptier with probability at least $1 - 2^{-\text{polylog}(n)}$. Whereas classically randomization gives the emptier a large advantage, in the variable-processor cup game the lower bound is the same!

1 Introduction

Definition and Motivation. The *cup game* is a multi-round game in which the two players, the *filler* and the *emptier*, take turns adding and removing water from cups. On each round of the classic *p*-processor cup game on n cups, the filler first distributes p units of water among the n cups with at most 1 unit to any particular cup (without this restriction the filler can trivially achieve unbounded backlog by placing all of its fill in a single cup every round), and then the emptier removes at most 1 unit of water from each of p cups.¹ The game has been studied for *adaptive* fillers, i.e. fillers that can observe the emptier’s actions, and for *oblivious* fillers, i.e. fillers that cannot observe the emptier’s actions.

The cup game naturally arises in the study of processor-scheduling. The incoming water added by the filler represents work added to the system at time steps. At each time step after the new work comes in, each of p processors must be allocated to a task which they will achieve 1 unit of progress on before the next time step. The assignment of processors to tasks is modeled by the emptier deciding which cups to empty from. The backlog of the system is the largest amount of work left on any given task; in the cup game the *backlog* of the cups is the fill of the fullest cup at a given state. In analyzing a cup game we aim to prove upper and lower bounds on backlog.

Previous Work. The bounds on backlog are well known for the case where $p = 1$, i.e. the *single-processor cup game*. In the single-processor cup game an adaptive filler can achieve backlog $\Omega(\log n)$ and a greedy emptier never lets backlog exceed $O(\log n)$. In the randomized version of the single-processor cup game, i.e. when the filler is oblivious, which can be interpreted as a smoothed analysis of the deterministic version, the emptier never lets backlog exceed $O(\log \log n)$, and a filler can achieve backlog $\Omega(\log \log n)$.

^{*}Supported by a Hertz fellowship and a NSF GRFP fellowship

[†]Supported by MIT

¹Note that negative fill is not allowed, so if the emptier empties from a cup with fill below 1 that cup’s fill becomes 0.

Recently Kuszmaul has established bounds on the case where $p > 1$, i.e. the **multi-processor cup game** [4]. Kuszmaul showed that a greedy emptier never lets backlog exceed $O(\log n)$. He also proved a lower bound of $\Omega(\log(n - p))$ on backlog. Recently we showed a lower bound of $\Omega(\log n - \log(n - p))$. Combined, these lower bounds imply a lower bound of $\Omega(\log n)$. Kuszmaul also established an upper bound of $O(\log \log n + \log p)$ against oblivious fillers, and a lower bound of $\Omega(\log \log n)$. Tight bounds on backlog against an oblivious filler are not yet known for large p .

The Variable-Processor Cup Game. We investigate a new variant of the classic p -processor cup game which we call the **variable-processor cup game**. In the variable-processor cup game the filler is allowed to change p (the total amount of water that the filler adds, and the emptier removes, from the cups per round) at the beginning of each round. Note that we do not allow the resources of the filler and emptier to vary separately; just like in the classic cup game we take the resources of the filler and emptier to be identical. This restriction is crucial; if the filler has more resources than the emptier, then the filler could trivially achieve unbounded backlog, as average fill will increase by at least some positive constant at each round. Taking the resources of the players to be identical makes the game balanced, and hence interesting.

The variable-processor cup game models the natural situation where many users are all on a server, and the number of processors allocated to each user is variable as other users get some portion of the processors.

A priori having variable resources offers neither player a clear advantage: lower values of p mean that the emptier is at more of a discretization disadvantage but also mean that the filler can “anchor” fewer cups.² Furthermore, at any fixed value of p upper bounds have been proven. For instance, regardless of p a greedy emptier prevents an adaptive filler from having backlog greater than $O(\log n)$. Switching between different values of p , all of which the filler cannot individually use to get backlog larger than $O(\log n)$ is not obviously going to help the filler achieve larger backlog. We hoped that the variable-processor cup game could be simulated in the classic multi-processor cup game, because the extra ability given to the filler does not seem very strong.

²A useful part of many filling algorithms is maintaining an “anchor” set of “anchored” cups. The filler always places 1 unit of water in each anchored cup. This ensures that the fill of an anchored cup never decreases after it is placed in the anchor set.

However, we show that attempts at simulating the variable-processor cup game are futile because the variable-processor cup game is vastly different from the classic multi-processor cup game.

Outline and Results. In Section 2 we establish the conventions and notations we will use to discuss the variable-processor cup game.

In Section 3 we provide an inductive proof of a lower bound on backlog with an adaptive filler. Theorem 1 states that a filler can achieve backlog $\Omega(n^{1-\varepsilon})$ for any constant $\varepsilon > 0$ in quasi-polynomial running time. Proposition 2 also provides an extremal strategy that achieves backlog $\Omega(n)$ in incredibly long games: it has $O(n!)$ running time.

In Section 4 we prove a novel invariant maintained by the greedy emptier. In particular Theorem 2 establishes that a greedy emptier keeps the average fill of the k fullest cups at most $2n - k$. In particular this implies (setting $k = 1$) that a greedy emptier prevents backlog from exceeding $O(n)$.

The lower bound and upper bound agree; our analysis is tight for adaptive fillers!

In Section 5 we prove a lower bound on backlog with an oblivious filler. Theorem 3 states that an oblivious filler can achieve backlog $\Omega(n^{1-\varepsilon})$ for any constant $\varepsilon > 0$ in quasi-polynomial time with probability at least $1 - 2^{-\text{polylog}(n)}$. Theorem 3 only applies to a certain class of emptiers: “greedy-like emptiers”. Nonetheless, this class of emptiers is very interesting; it contains the emptiers that are used in upper bound proofs. It is shocking that randomization doesn’t help the emptier in this game; being oblivious seems like a large disadvantage for the filler!

2 Preliminaries

The cup game consists of a sequence of rounds. On the t -th round, the state starts as S_t . The filler chooses the number of processors p_t for the round. Then the filler distributes p_t units of water among the cups (with at most 1 unit of water to any particular cup). After this the game is in an intermediate state on round t , which we call state I_t . Then the emptier chooses p_t cups to empty at most 1 unit of water from. Note that if the fill of a cup that the emptier empties from is less than 1 the emptier reduces the fill of this cup to 0 by emptying from it; we say that the emptier **zeroes out** a cup at round t if the emptier empties, on round t , from a cup with fill at state I_t that is less than 1. Note that on any round where the emptier zeroes out a cup the emptier has removed less fill than the filler has added; hence the average fill will increase. This concludes the round;

the state of the game is now S_{t+1} .

Denote the fill of a cup c by $\text{fill}(c)$. Let the **mass** of a set of cups X be $m(X) = \sum_{c \in X} \text{fill}(c)$. Denote the average fill of a set of cups X by $\mu(X)$. Note that $\mu(X)|X| = m(X)$.

Let the **rank** of a cup at a given state be its position in a list of the cups sorted by fill at the given state, breaking ties arbitrarily but consistently. For example, the fullest cup at a state has rank 1, and the least full cup has rank n . Let $[n] = \{1, 2, \dots, n\}$, let $i + [n] = \{i + 1, i + 2, \dots, i + n\}$.

Many of our lower bound proofs will adopt the convention of allowing for negative fill. We call this the **negative-fill cup game**. Specifically, in the negative-fill cup game, when the emptier empties from a cup its fill always decreases by exactly 1: there is no zeroing out. Negative-fill can be interpreted as fill below some average fill. Measuring fill like this is important however, as our lower bound results are used recursively, building on the average fill already achieved. Note that it is strictly easier for the filler to achieve high backlog when cups can zero out, because then some of the emptier's resources are wasted. On the other hand, during the upper bound proof we show that a greedy emptier maintains the desired invariants even if cups zero out. This is crucial as the game is harder for the emptier when cups can zero out. Some results are proved for the variable-processor negative-fill cup game, and some results are proved for the single-processor negative-fill cup game.

3 Adaptive Filler Lower Bound

In this section we give a $2^{\text{polylog } n}$ -time filling strategy that achieves backlog $n^{1-\epsilon}$ for any positive constant ϵ . We also give a $O(n!)$ -time filling strategy that achieves backlog $\Omega(n)$.

We begin with a simple proposition that gives backlog $1/2$ for two cups.

Proposition 1. *Consider an instance of the negative-fill 1-processor cup game on 2 cups, and let the cups start in any state where the average fill is 0. There is an $O(1)$ -step adaptive filling strategy that achieves backlog at least $1/2$.*

Proof. Let the fills of the 2 cups start as x and $-x$ for some $x \geq 0$. If $x \geq 1/2$ the algorithm need not do anything. Otherwise, the filling strategy adds $1/2 - x$ fill to the cup with fill x , and adds $1/2 + x$ fill to the cup with fill $-x$. This results in 2 cups both having fill $1/2$; the emptier then empties from one of these, and leaves a cup with fill $1/2$, as desired. \square

Next we prove the **Amplification Lemma**.

Lemma 1 (Adaptive Amplification Lemma). *Let $\delta \in (0, 1/2)$ be a parameter. Let $\text{alg}(f)$ be an adaptive filling strategy that achieves backlog $f(n) < n$ in the negative-fill variable-processor cup game on n cups in running time $T(n)$ starting from any initial cup state where the average fill is 0.*

Then there exists an adaptive filling strategy $\text{alg}(f')$ that achieves backlog $f'(n)$ satisfying

$$f'(n) \geq (1 - \delta)f(\lfloor (1 - \delta)n \rfloor) + f(\lceil \delta n \rceil)$$

and $f'(n) \geq f(n)$ in the negative-fill variable-processor cup game on n cups in running time

$$T'(n) \leq n^2 \delta \cdot T(\lfloor (1 - \delta)n \rfloor) + T(\lceil \delta n \rceil)$$

starting from any initial cup state where the average fill is 0.

Before proving the Amplification Lemma, we briefly motivate it. We call $\text{alg}(f')$, the filling strategy created by the Amplification Lemma, the **amplification** of $\text{alg}(f)$. As suggested by the name, $\text{alg}(f')$ will be able to achieve higher backlog than $\text{alg}(f)$. In particular, we will show that by starting with a filling strategy $\text{alg}(f_0)$ for achieving constant backlog and then recursively forming a sufficiently long sequence of filling strategies $\text{alg}(f_0), \text{alg}(f_1), \dots, \text{alg}(f_{i_*})$ with $\text{alg}(f_{i+1})$ the amplification of $\text{alg}(f_i)$, we eventually get a filling strategy for achieving $\text{poly}(n)$ backlog.

Proof of Amplification Lemma. The filler defaults to using $\text{alg}(f)$ if $f(n) \geq (1 - \delta)f(\lfloor (1 - \delta)n \rfloor) + f(\lceil \delta n \rceil)$; in this case using $\text{alg}(f)$ achieves the desired backlog in the desired running time. In the rest of the proof, we describe our strategy for achieving backlog $(1 - \delta)f(\lfloor (1 - \delta)n \rfloor) + f(\lceil \delta n \rceil)$.

Let A , the **anchor set**, be initialized to consist of the $\lceil n\delta \rceil$ fullest cups, and let B the **non-anchor set** be initialized to consist of the rest of the cups (so $|B| = \lfloor (1 - \delta)n \rfloor$). Let $h = (1 - \delta)f(\lfloor (1 - \delta)n \rfloor)$.

The filler's strategy is roughly as follows:

Step 1: Get $\mu(A) \geq h$ by using $\text{alg}(f)$ repeatedly on B to achieve cups with fill at least $\mu(B) + f(|B|)$ in B and then swapping these into A .

Step 2: Use $\text{alg}(f)$ once on A to obtain some cup with fill $\mu(A) + f(|A|)$.

Note that in order to use $\text{alg}(f)$ on subsets of the cups the filler will need to vary p .

We now describe how to achieve Step 1, which is complicated by the fact that the emptier may attempt to prevent the filler from achieving high fill in a cup in B .

The filling strategy always places 1 unit of water in each anchor cup. This ensures that no cups in the anchor set ever have their fill decrease. If the

emptier wishes to keep the average fill of the anchor cups from increasing, then emptier must empty from every anchor cup on each step. If the emptier fails to do this on a given round, then we say that the emptier has **neglected** the anchor cups.

We say that the filler **applies** $\text{alg}(f)$ to B if it follows the filling strategy $\text{alg}(f)$ on B while placing 1 unit of water in each anchor cup. An application of $\text{alg}(f)$ to B is said to be **successful** if A is never neglected during the application of $\text{alg}(f)$ to B . The filler uses a procedure that we call a **swapping-process** to achieve the desired average fill in A . In a swapping-process, the filler repeatedly applies $\text{alg}(f)$ to B until a successful application occurs, and then takes the cup generated by $\text{alg}(f)$ within B on this successful application with fill at least $\mu(B) + f(|B|)$ and swaps it with the least full cup in A . If the average fill in A ever reaches h , then the algorithm immediately halts (even if it is in the middle of a swapping-process) and is complete.

Note that

$$\mu(A) \cdot |A| + \mu(B) \cdot |B| = 0,$$

so

$$\mu(A) = -\mu(B) \cdot \frac{\lfloor (1-\delta)n \rfloor}{\lceil \delta n \rceil} \geq -\frac{1-\delta}{\delta} \mu(B).$$

Thus, if at any point B has average fill lower than $-h \cdot \delta / (1-\delta)$, then A has average fill at least h , so the algorithm is finished. Thus we can assume in our analysis that

$$\mu(B) \geq -h \cdot \delta / (1-\delta). \quad (1)$$

We will now show that during each swapping process, the filler applies $\text{alg}(f)$ to B at most $hn\delta + 1$ times. Each time the emptier neglects the anchor set, the mass of the anchor set increases by 1. If the emptier neglects the anchor set $hn\delta + 1$ times, then the average fill in the anchor set increases by more than h , so the desired average fill is achieved in the anchor set. Thus the swapping process consists of at most $hn\delta + 1$ applications of $\text{alg}(f)$.

Consider the fill of a cup c swapped into A at the end of a swapping-process. Cup c 's fill is at least $\mu(B) + f(|B|)$, which by (1) is at least

$$-h \cdot \frac{\delta}{1-\delta} + f(\lfloor n(1-\delta) \rfloor) = (1-\delta)f(\lfloor n(1-\delta) \rfloor) = h.$$

Thus the algorithm for Step 1 succeeds within $|A| = \lceil \delta n \rceil$ swapping-processes, since at the end of the $|A|$ -th swapping process every cup in A has fill at least h , or the algorithm halted before $|A|$ swapping-processes because it already achieved $\mu(A) \geq h$.

Now the filler performs Step 2, i.e. the filler applies $\text{alg}(f)$ to A , and hence achieves a cup with fill at least

$$\mu(A) + f(|A|) \geq (1-\delta)f(\lfloor (1-\delta)n \rfloor) + f(\lceil \delta n \rceil),$$

as desired.

Now we analyze the running time of the filling strategy $\text{alg}(f')$. First, recall that in Step 1 $\text{alg}(f')$ calls $\text{alg}(f)$ on a set of size $\lfloor (1-\delta)n \rfloor$ as many as $hn\delta + 1$ times. Because we mandate that $h < n$, Step 1 contributes no more than $(n \cdot n\delta) \cdot T(|B|)$ to the running time. Step 2 requires applying $\text{alg}(f)$ to $|A|$ cups one time, and hence contributes $T(|A|)$ to the running time. Summing these we have

$$T'(n) \leq n^2\delta \cdot T(\lfloor (1-\delta)n \rfloor) + T(\lceil \delta n \rceil).$$

□

We next show that by recursively using the Amplification Lemma we can achieve backlog $n^{1-\varepsilon}$.

Theorem 1. *There is an adaptive filling strategy for the variable-processor cup game on n cups that achieves backlog $\Omega(n^{1-\varepsilon})$ for any constant $\varepsilon > 0$ of our choice in running time $2^{O(\log^2 n)}$.*

Proof. Take constant $\varepsilon \in (0, 1/2)$. Let c, δ be parameters, with $c \in (0, 1)$, $0 < \delta \ll 1/2$, these will be chosen later as functions of ε . We show how to achieve backlog at least $cn^{1-\varepsilon} - 1$.

By Proposition 1 there exists a constant n_0 such that a filler can achieve backlog 1 on n_0 cups (e.g., $n_0 = 1000$ works). Let $\text{alg}(f_0)$ be the filling strategy described in Proposition 1, where $f_0(k) \geq 1$ for all $k \geq n_0$.

Next, using the Amplification Lemma we recursively construct $\text{alg}(f_{i+1})$ as the amplification of $\text{alg}(f_i)$ for $i \geq 0$.

Define a sequence g_i with

$$g_i = \begin{cases} \lceil 16/\delta \rceil, & i = 0, \\ \lfloor g_{i-1}/(1-\delta) \rfloor & i \geq 1 \end{cases}$$

We claim the following regarding this construction:

Claim 1. *For all $i \geq 0$,*

$$f_i(k) \geq ck^{1-\varepsilon} - 1 \text{ for all } k \in [g_i]. \quad (2)$$

Proof. We prove Claim 1 by induction on i . For $i = 0$, the base case, (2) can be made true by taking c and δ sufficiently small. In particular, we choose $c = \Theta(1)$ small enough to make $cn_0^{1-\varepsilon} - 1 \leq 0$, which implies (2) holds for $k \in [n_0]$ by monotonicity of $ck^{1-\varepsilon} - 1$; we also choose δ small enough to make $g_0 \geq n_0$, and we choose c small enough to make

$cg_0^{1-\varepsilon} - 1 \leq f_0(g_0) = 1$, which implies (2) holds for $k \in [n_0, g_0]$ by monotonicity of $ck^{1-\varepsilon} - 1$.³

As our inductive hypothesis we assume (2) for f_i ; we aim to show that (2) holds for f_{i+1} . Note that, by design of g_i , if $k \leq g_{i+1}$ then $\lfloor k \cdot (1 - \delta) \rfloor \leq g_i$. Consider any $k \in [g_{i+1}]$. First we deal with the trivial case where $k \leq g_0$. In this case

$$f_{i+1}(k) \geq f_i(k) \geq \dots \geq f_0(k) \geq ck^{1-\varepsilon} - 1.$$

Now we consider the case where $k \geq g_0$. Since f_{i+1} is the amplification of f_i we have

$$f_{i+1}(k) \geq (1 - \delta)f_i(\lfloor (1 - \delta)k \rfloor) + f_i(\lceil \delta k \rceil).$$

By our inductive hypothesis, which applies as $\lceil \delta k \rceil \leq g_i$, $\lfloor k \cdot (1 - \delta) \rfloor \leq g_i$, we have

$$f_{i+1}(k) \geq (1 - \delta)(c \cdot \lfloor (1 - \delta)k \rfloor^{1-\varepsilon} - 1) + c \lceil \delta k \rceil^{1-\varepsilon} - 1.$$

Dropping the floor and ceiling, incurring a -1 for dropping the floor, we have

$$f_{i+1}(k) \geq (1 - \delta)(c \cdot ((1 - \delta)k - 1)^{1-\varepsilon} - 1) + c(\delta k)^{1-\varepsilon} - 1.$$

Because $(x - 1)^{1-\varepsilon} \geq x^{1-\varepsilon} - 1$, due to the fact that $x \mapsto x^{1-\varepsilon}$ is a sub-linear sub-additive function, we have

$$f_{i+1}(k) \geq (1 - \delta)c \cdot (((1 - \delta)k)^{1-\varepsilon} - 2) + c(\delta k)^{1-\varepsilon} - 1.$$

Moving the $ck^{1-\varepsilon}$ to the front we have

$$f_{i+1}(k) \geq ck^{1-\varepsilon} \cdot \left((1 - \delta)^{2-\varepsilon} + \delta^{1-\varepsilon} - \frac{2(1 - \delta)}{k^{1-\varepsilon}} \right) - 1.$$

Because $(1 - \delta)^{2-\varepsilon} \geq 1 - (2 - \varepsilon)\delta$, a fact called Bernoulli's Identity, we have

$$f_{i+1}(k) \geq ck^{1-\varepsilon} \cdot \left(1 - (2 - \varepsilon)\delta + \delta^{1-\varepsilon} - \frac{2(1 - \delta)}{k^{1-\varepsilon}} \right) - 1.$$

Of course $-2(1 - \delta) \geq -2$, so

$$f_{i+1}(k) \geq ck^{1-\varepsilon} \cdot \left(1 - (2 - \varepsilon)\delta + \delta^{1-\varepsilon} - \frac{2}{k^{1-\varepsilon}} \right) - 1.$$

Because

$$\frac{-2}{k^{1-\varepsilon}} \geq \frac{-2}{g_0^{1-\varepsilon}} \geq -2(\delta/16)^{1-\varepsilon} \geq -\delta^{1-\varepsilon}/2,$$

which follows from our choice of $g_0 = \lceil 16/\delta \rceil$ and the restriction $\varepsilon < 1/2$, we have

$$f_{i+1}(k) \geq ck^{1-\varepsilon} \cdot (1 - (2 - \varepsilon)\delta + \delta^{1-\varepsilon} - \delta^{1-\varepsilon}/2) - 1.$$

³Note that it is important here that ε and δ are constants, that way c is also a constant.

Finally, combining terms we have

$$f_{i+1}(k) \geq ck^{1-\varepsilon} \cdot (1 - (2 - \varepsilon)\delta + \delta^{1-\varepsilon}/2) - 1.$$

Because $\delta^{1-\varepsilon}$ dominates δ for sufficiently small δ , there is a choice of $\delta = \Theta(1)$ such that

$$1 - (2 - \varepsilon)\delta + \delta^{1-\varepsilon}/2 \geq 1.$$

Taking δ to be this small we have,

$$f_{i+1}(k) \geq ck^{1-\varepsilon} - 1,$$

completing the proof. We remark that the choices of c, δ are the same for every i in the inductive proof, and depend only on ε . \square

To complete the proof, we will show that g_i grows exponentially in i . Thus, after there exists $i_* \leq O(\log n)$ such that $g_{i_*} \geq n$, and hence we have an algorithm $\text{alg}(f_{i_*})$ that achieves backlog $cn^{1-\varepsilon} - 1$ on n cups, as desired.

We lower bound the sequence g_i with another sequence g'_i defined as

$$g'_i = \begin{cases} 4/\delta, & i = 0 \\ g'_{i-1}/(1 - \delta) - 1, & i > 0. \end{cases}$$

Solving this recurrence, we find

$$g'_i = \frac{4 - (1 - \delta)^2}{\delta} \frac{1}{(1 - \delta)^i} \geq \frac{1}{(1 - \delta)^i},$$

which clearly exhibits exponential growth. In particular, let $i_* = \lceil \log_{1/(1-\delta)} n \rceil$. Then,

$$g_{i_*} \geq g'_{i_*} \geq n,$$

as desired.

Let the running time of $f_i(n)$ be $T_i(n)$. From the Amplification Lemma we have following recurrence bounding $T_i(n)$:

$$\begin{aligned} T_i(n) &\leq n^2\delta \cdot T_{i-1}(\lfloor (1 - \delta)n \rfloor) + T_{i-1}(\lceil \delta n \rceil) \\ &\leq 2n^2T_{i-1}(\lfloor (1 - \delta)n \rfloor). \end{aligned}$$

It follows that $\text{alg}(f_{i_*})$, recalling that $i_* \leq O(\log n)$, has running time

$$T_{i_*}(n) \leq (2n^2)^{O(\log n)} \leq 2^{O(\log^2 n)}$$

as desired. \square

Now we provide a very simple construction that can achieve backlog $\Omega(n)$ in very long games. The construction can be interpreted as the same argument

as in Theorem 1 but with an extremal setting of δ to $\Theta(1/n)$.⁴

Proposition 2. *There is an adaptive filling strategy that achieves backlog $\Omega(n)$ in time $O(n!)$.*

Proof. We start, as in the proof of Theorem 1, with an algorithm $\text{alg}(f_0)$ for achieving backlog $f_0(k) \geq 1$ on $k \geq n_0$ cups, which is possible by Proposition 1. For $i > 0$ we construct $\text{alg}(f_i)$ as the amplification of $\text{alg}(f_{i-1})$ using the Amplification Lemma with parameter $\delta = 1/(i+1)$.

We claim the following regarding this construction:

Claim 2. *For all $i \geq 0$,*

$$f_i((i+1) \cdot n_0) \geq \sum_{j=0}^i \left(1 - \frac{j}{i+1}\right). \quad (3)$$

Proof. We prove Claim 2 by induction on i . When $i = 0$, the base case, (3) becomes $f_0(n_0) \geq 1$ which is true. Assuming (3) for f_{i-1} , we now show (3) holds for f_i . Because f_i is the amplification of f_{i-1} with $\delta = 1/(i+1)$, we have by the Amplification Lemma

$$f_i((i+1) \cdot n_0) \geq \left(1 - \frac{1}{i+1}\right) f_{i-1}(i \cdot n_0) + f_{i-1}(n_0).$$

Since $f_{i-1}(n_0) \geq f_0(n_0) \geq 1$ we have

$$f_i((i+1) \cdot n_0) \geq \left(1 - \frac{1}{i+1}\right) f_{i-1}(i \cdot n_0) + 1.$$

Using the inductive hypothesis we have

$$f_i((i+1) \cdot n_0) \geq \left(1 - \frac{1}{i+1}\right) \sum_{j=0}^{i-1} \left(1 - \frac{j}{i}\right) + 1.$$

Note that

$$\begin{aligned} \left(1 - \frac{1}{i+1}\right) \cdot \left(1 - \frac{j}{i}\right) &= \frac{i}{i+1} \cdot \frac{i-j}{i} \\ &= \frac{i-j}{i+1} \\ &= 1 - \frac{j+1}{i+1}. \end{aligned}$$

Thus we have

$$f_i((i+1) \cdot n_0) \geq \sum_{j=1}^i \left(1 - \frac{j}{i+1}\right) + 1 = \sum_{j=0}^i \left(1 - \frac{j}{i+1}\right),$$

as desired. \square

⁴Or more precisely, setting δ in each level of recursion to be $\Theta(1/n)$, where n is the subproblem size; note in particular that δ changes between levels of recursion, which was not the case in the proof of Theorem 1.

Let $i_* = \lfloor n/n_0 \rfloor - 1$, which by design satisfies $(i_* + 1)n_0 \leq n$. By Claim 2 we have

$$f_{i_*}((i_* + 1) \cdot n_0) \geq \sum_{j=0}^{i_*} \left(1 - \frac{j}{i_* + 1}\right) = i_*/2 + 1.$$

As $i_* = \Theta(n)$, we have thus shown that $\text{alg}(f_{i_*})$ can achieve backlog $\Omega(n)$ on n cups.

Let T_i be the running time of $\text{alg}(f_i)$. The recurrence for the running time of f_{i_*} is

$$T_i(n) \leq n \cdot n_0 T_{i-1}(n - n_0) + O(1).$$

Clearly $T_{i_*}(n) \leq O(n!)$. \square

4 Upper Bound

In this section we analyze the *greedy emptier*, which always empties from the p fullest cups. We prove in Corollary 1 that the greedy emptier prevents backlog from exceeding $O(n)$.

In order to analyze the greedy emptier, we establish a system of invariants that hold at every step of the game.

Let $\mu_S(X)$ and $m_S(X)$ denote the average fill and the mass, respectively, of a set of cups X at state S (e.g. $S = S_t$ or $S = I_t$).⁵ Let $S(\{r_1, \dots, r_m\})$ denote the set of cups of ranks r_1, r_2, \dots, r_m at state S . We will use concatenation of sets to denote unions, i.e. $AB = A \cup B$.

The main result of the section is the following theorem.

Theorem 2. *In the variable-processor cup game on n cups, the greedy emptier maintains, at every step t , the invariants*

$$\mu_{S_t}(S_t([k])) \leq 2n - k \quad (4)$$

for all $k \in [n]$.

By applying Theorem 2 to the case of $k = 1$, we arrive at a bound on backlog:

Corollary 1. *In the variable-processor cup game on n cups, the greedy emptying strategy never lets backlog exceed $O(n)$.*

⁵Note that in the lower bound proofs (i.e. Section 3 and Section 5) when we use the notation m (for mass) and μ (for average fill), we omit the subscript indicating the state at which the properties are measured. In those proofs the state is implicitly clear. However, in this section it will be useful to make the state S explicit in the notation.

Proof of Theorem 2. We prove the invariants by induction on t . The invariants hold trivially for $t = 1$ (the base case for the inductive proof): the cups start empty so $\mu_{S_1}(S_1([k])) = 0 \leq 2n - k$ for all $k \in [n]$.

Fix a round $t \geq 1$, and any $k \in [n]$. We assume the invariants for all values of $k' \in [n]$ for state S_t (we will only explicitly use two of the invariants for each k , but the invariants that we need depend on the choice of p_t by the filler) and show that the invariant on the k fullest cups holds on round $t + 1$, i.e. that

$$\mu_{S_{t+1}}(S_{t+1}([k])) \leq 2n - k.$$

Note that because the emptier is greedy it always empties from the cups $I_t([p_t])$. Let A , with $a = |A|$, be $A = I_t([\min(k, p_t)]) \cap S_{t+1}([k])$; A consists of the cups that are among the k fullest cups in I_t , are emptied from, and are among the k fullest cups in S_{t+1} . Let B , with $b = |B|$, be $I_t([\min(k, p_t)]) \setminus A$; B consists of the cups that are among the k fullest cups in state I_t , are emptied from, and are not among the k fullest cups in S_{t+1} . Let $C = I_t(a + b + [k - a])$, with $c = k - a = |C|$; C consists of the cups with ranks $a + b + 1, \dots, k + b$ in state I_t . The set C is defined so that $S_{t+1}([k]) = AC$, since once the cups in B are emptied from, the cups in B are not among the k fullest cups, so cups in C take their places among the k fullest cups.

Note that $k - a \geq 0$ as $a + b \leq k$, and also $|ABC| = k + b \leq n$, because by definition the b cups in B must not be among the k fullest cups in state S_{t+1} so there are at least $k + b$ cups. Note that $a + b = \min(k, p_t)$. We also have that $A = I_t([a])$ and $B = I_t(a + [b])$, as every cup in A must have higher fill than all cups in B in order to remain above the cups in B after 1 unit of water is removed from all cups in AB .

We now establish the following claim, which we call the *interchangeability of cups*:

Claim 3. *There exists a cup state S'_t such that: (a) S'_t satisfies the invariants (4), (b) $S'_t(r) = I_t(r)$ for all ranks $r \in [n]$, and (c) the filler can legally place water into cups in order to transform S'_t into I_t .*

Proof. Fix $r \in [n]$. We will show that S_t can be transformed into a state S'_t by relabelling only cups with ranks in $[r]$ such that (a) S'_t satisfies the invariants (4), (b) $S'_t([r]) = I_t([r])$ and (c) the filler can legally place water into cups in order to transform S'_t into I_t .

Say there are cups x, y with $x \in S_t([r]) \setminus I_t([r])$, $y \in I_t([r]) \setminus S_t([r])$. Let the fills of cups x, y at state S_t be f_x, f_y ; note that

$$f_x > f_y. \quad (5)$$

Let the amount of fill that the filler adds to these cups be $\Delta_x, \Delta_y \in [0, 1]$; note that

$$f_x + \Delta_x < f_y + \Delta_y. \quad (6)$$

Define a new state S'_t where cup x has fill f_y and cup y has fill f_x . Note that the filler can transform state S'_t into state I_t by placing water into cups as before, except changing the amount of water placed into cups x and y to be $f_x - f_y + \Delta_x$ and $f_y - f_x + \Delta_y$, respectively.

In order to verify that the transformation from S'_t to I_t is a valid step for the filler, one must check three conditions. First, the amount of water placed by the filler is unchanged: this is because $(f_x - f_y + \Delta_x) + (f_y - f_x + \Delta_y) = \Delta_x + \Delta_y$. Second, the fills placed in cups x and y are at most 1: this is because $f_x - f_y + \Delta_x < \Delta_y \leq 1$ (by (6)) and $f_y - f_x + \Delta_x < \Delta_x \leq 1$ (by (5)). Third, the fills placed in cups x and y are non-negative: this is because $f_x - f_y + \Delta_x > \Delta_x \geq 0$ (by (5)) and $f_y - f_x + \Delta_y > \Delta_y \geq 0$ (by (6)).

We can repeatedly apply this process to swap each cup in $I_t([r]) \setminus S_t([r])$ into being in $S'_t([r])$. At the end of this process we will have some state S'_t for which $S'_t([r]) = I_t([r])$. Note that S'_t is simply a relabeling of S_t , hence it must satisfy the same invariants (4) satisfied by S_t . Further, S'_t can be transformed into I_t by a valid filling step.

Now we repeatedly apply this process, in descending order of ranks. In particular, we have the following process: create a sequence of states by starting with S_t^{n-1} , and to get to state S'_t from state S_t^{r+1} apply the process described above. Note that S_t^{n-1} satisfies $S_t^{n-1}([n-1]) = I_t([n-1])$ and thus also $S_t^{n-1}(n) = I_t(n)$. If S_t^{r+1} satisfies $S_t^{r+1}(r') = I_t(r')$ for all $r' > r+1$ then S'_t satisfies $S'_t(r') = I_t(r')$ for all $r > r$, because the transition from S_t^{r+1} to S'_t has not changed the labels of any cups with ranks in $(r+1, n]$, but the transition does enforce $S'_t([r]) = I_t([r])$, and consequently $S'_t(r+1) = I_t(r+1)$. We continue with the sequential process until arriving at state S'_t in which we have $S'_t(r) = I_t(r)$ for all r . Throughout the process each S'_t has satisfied the invariants (4), so S'_t satisfies the invariants (4). Further, throughout the process from each S'_t it is possible to legally place water into cups in order to transform S'_t into I_t .

Hence S'_t satisfies all the properties desired, and the proof of Claim 3 is complete. \square

Claim 3 tells us that we may assume without loss of generality that $S_t(r) = I_t(r)$ for each rank $r \in [n]$. We will make this assumption for the rest of the proof.

In order to complete the proof of the theorem, we break it into three cases.

Claim 4. *If some cup in A zeroes out in round t , then the invariant $\mu_{S_{t+1}}(S_{t+1}([k])) \leq 2n - k$ holds.*

Proof. Say a cup in A zeroes out in step t . Of course

$$m_{S_{t+1}}(I_t([a-1])) \leq (a-1)(2n - (a-1))$$

because the $a-1$ fullest cups must have satisfied the invariant (with $k = a-1$) on round t . Moreover, because $\text{fill}_{S_{t+1}}(I_{t+1}(a)) = 0$

$$m_{S_{t+1}}(I_t([a])) = m_{S_{t+1}}(I_t([a-1])).$$

Combining the above equations, we get that

$$m_{S_{t+1}}(A) \leq (a-1)(2n - (a-1)).$$

Furthermore, the fill of all cups in C must be at most 1 at state I_t to be less than the fill of the cup in A that zeroed out. Thus,

$$\begin{aligned} m_{S_{t+1}}(S_{t+1}([k])) &= m_{S_{t+1}}(AC) \\ &\leq (a-1)(2n - (a-1)) + k - a \\ &= a(2n - a) + a - 2n + a - 1 + k - a \\ &= a(2n - a) + (k - n) + (a - n) - 1 \\ &< a(2n - a) \end{aligned}$$

as desired. As k increases from 1 to n , $k(2n - k)$ strictly increases (it is a quadratic in k that achieves its maximum value at $k = n$). Thus $a(2n - a) \leq k(2n - k)$ because $a \leq k$. Therefore,

$$m_{S_{t+1}}(S_{t+1}([k])) \leq k(2n - k).$$

□

Claim 5. *If no cups in A zero out in round t and $b = 0$, then the invariant $\mu_{S_{t+1}}(S_{t+1}([k])) \leq 2n - k$ holds.*

Proof. If $b = 0$, then $S_{t+1}([k]) = S_t([k])$. During round t the emptier removes a units of fill from the cups in $S_t([k])$, specifically the cups in A . The filler cannot have added more than k fill to these cups, because it can add at most 1 fill to any given cup. Also, the filler cannot have added more than p_t fill to the cups because this is the total amount of fill that the filler is allowed to add. Hence the filler adds at most $\min(p_t, k) = a + b = a + 0 = a$ fill to these cups. Thus the invariant holds:

$$m_{S_{t+1}}(S_{t+1}([k])) \leq m_{S_t}(S_t([k])) + a - a \leq k(2n - k).$$

□

The remaining case, in which no cups in A zero out and $b > 0$ is the most technically interesting.

Claim 6. *If no cups in A zero out on round t and $b > 0$, then the invariant $\mu_{S_{t+1}}(S_{t+1}([k])) \leq 2n - k$ holds.*

Proof. Because $b > 0$ and $a + b \leq k$ we have that $a < k$, and $c = k - a > 0$. Recall that $S_{t+1}([k]) = AC$, so the mass of the k fullest cups at S_{t+1} is the mass of AC at S_t plus any water added to cups in AC by the filler, minus any water removed from cups in AC by the emptier. The emptier removes exactly a units of water from AC . The filler adds no more than p_t units of water to AC (because the filler adds at most p_t total units of water per round) and the filler also adds no more than $k = |AC|$ units of water to AC (because the filler adds at most 1 unit of water to each of the k cups in AC). Thus, the filler adds no more than $a + b = \min(p_t, k)$ units of water to AC . Combining these observations we have:

$$m_{S_{t+1}}(S_{t+1}([k])) \leq m_{S_t}(AC) + b. \quad (7)$$

The key insight necessary to bound this is to notice that larger values for $m_{S_t}(A)$ correspond to smaller values for $m_{S_t}(C)$ because of the invariants; the higher fill in A **pushes down** the fill that C can have. By capturing the pushing-down relationship combinatorially we will achieve the desired inequality.

We can upper bound $m_{S_t}(C)$ by

$$\begin{aligned} m_{S_t}(C) &\leq \frac{c}{b+c} m_{S_t}(BC) \\ &= \frac{c}{b+c} (m_{S_t}(ABC) - m_{S_t}(A)) \end{aligned}$$

because $\mu_{S_t}(C) \leq \mu_{S_t}(B)$ without loss of generality by the interchangeability of cups. Thus we have

$$m_{S_t}(AC) \leq m_{S_t}(A) + \frac{c}{b+c} m_{S_t}(BC) \quad (8)$$

$$= \frac{c}{b+c} m_{S_t}(ABC) + \frac{b}{b+c} m_{S_t}(A). \quad (9)$$

Note that the expression in (9) is monotonically increasing in both $\mu_{S_t}(ABC)$ and $\mu_{S_t}(A)$. Thus, by numerically replacing both average fills with their extremal values, $2n - |ABC|$ and $2n - |A|$. At this point the claim can be verified by straightforward (but quite messy) algebra (and by combining (7) with (9)). We instead give a more intuitive argument, in which we examine the right side of (8) combinatorially.

Consider a new configuration of fills F achieved by starting with state S_t , and moving water from BC into A until $\mu_F(A) = 2n - |A|$.⁶ This transformation

⁶Note that whether or not F satisfies the invariants is irrelevant.

increases (strictly increases if and only if we move a non-zero amount of water) the right side of (8). In particular, if mass $\Delta \geq 0$ fill is moved from BC to A , then the right side of (8) increases by $\frac{b}{b+c}\Delta \geq 0$. Note that the fact that moving water from BC into A increases the right side of (8) formally captures the way the system of invariants being proven forces a tradeoff between the fill in A and the fill in BC —that is, higher fill in A pushes down the fill that BC (and consequently C) can have.

Since $\mu_F(A)$ is above $\mu_F(ABC)$, the greater than average fill of A must be counter-balanced by the lower than average fill of BC . In particular we must have

$$(\mu_F(A) - \mu_F(ABC))|A| = (\mu_F(ABC) - \mu_F(BC))|BC|.$$

Note that

$$\begin{aligned} & \mu_F(A) - \mu_F(ABC) \\ &= (2n - |A|) - \mu_F(ABC) \\ &\geq (2n - |A|) - (2n - |ABC|) \\ &= |BC|. \end{aligned}$$

Hence we must have

$$\mu_F(ABC) - \mu_F(BC) \geq |A|.$$

Thus

$$\mu_F(BC) \leq \mu_F(ABC) - |A| \leq 2n - |ABC| - |A|. \quad (10)$$

Combing (8) with the fact that the transformation from S_t to F only increases the right side of (8), along with (10), we have the following bound:

$$\begin{aligned} m_{S_t}(AC) &\leq m_F(A) + c\mu_F(BC) \\ &\leq a(2n - a) + c(2n - |ABC| - a) \\ &\leq (a + c)(2n - a) - c(a + c + b) \\ &\leq (a + c)(2n - a - c) - cb. \end{aligned} \quad (11)$$

By (7) and (11), we have that

$$\begin{aligned} m_{S_{t+1}}(S_{t+1}([k])) &\leq m_{S_t}(AC) + b \\ &\leq (a + c)(2n - a - c) - cb + b \\ &= k(2n - k) - cb + b \\ &\leq k(2n - k), \end{aligned}$$

where the final inequality uses the fact that $c \geq 1$. This completes the proof of the claim. \square

We have shown the invariant holds for arbitrary k , so given that the invariants all hold at state S_t they also must all hold at state S_{t+1} . Thus, by induction we have the invariant for all rounds $t \in \mathbb{N}$. \square

5 Oblivious Filler Lower Bound

In this section we prove that, surprisingly, an oblivious filler can achieve backlog $n^{1-\varepsilon}$, although only against a certain class of “greedy-like” emptiers.

The **fill-range** of a set of cups at a state S is $\max_c \text{fill}_S(c) - \min_c \text{fill}_S(c)$. We call a cup configuration ***R-flat*** if the fill-range of the cups less than or equal to R ; note that in an R -flat cup configuration with average fill 0 all cups have fills in $[-R, R]$. We say an emptier is ***Δ -greedy-like*** if, whenever there are two cups with fills that differ by at least Δ , the emptier never empties from the less full cup without also emptying from the more full cup. That is, if on some round t , there are cups c_1, c_2 with $\text{fill}_{I_t}(c_1) > \text{fill}_{I_t}(c_2) + \Delta$, then a Δ -greedy-like emptier doesn’t empty from c_2 on round t unless it also empties from c_1 on round t . Note that a perfectly greedy emptier is 0-greedy-like. We call an emptier ***greedy-like*** if it is Δ -greedy-like for $\Delta \leq O(1)$.

With an oblivious filler we are only able to prove lower bounds on backlog against greedy-like emptiers; whether or not our results can be extended to a more general class of emptiers is an interesting open question. Nonetheless, greedy-like emptiers are of great interest because all the known randomized algorithms for the cup game are greedy-like [1, 4].

As a tool in our analysis we define a new variant of the cup game: In the p -processor ***E-extra-emptings S-skip-emptings*** negative-fill cup game on n cups, the filler distributes p units of water amongst the cups, and then the emptier empties from p or more, or less cups. In particular the emptier is allowed to do E extra emptings—we say that the emptier does an extra emptying if it empties from a cup beyond the p cups it typically is allowed to empty from—and is also allowed to skip S emptings—we say that the emptier skips an emptying if it doesn’t do an emptying—over the course of the game. Note that the emptier still cannot empty from the same cup twice on a single round. Also note that the emptier is allowed to skip extra emptings in addition to regular emptings. Also note that a Δ -greedy-like emptier must take into account extra emptings and skip emptings to determine valid moves.

For a Δ -greedy-like emptier let $R_\Delta = 2(2 + \Delta)$; we now prove a key property of these emptiers: there is an oblivious filling strategy, which we term ***fatalg***, that attains an R_Δ -flat cup configuration against a Δ -greedy-like emptier, given cups of a known starting fill-range.

Lemma 2. *Consider an R -flat cup configuration in the p -processor E -extra-emptings S -skip-emptings*

negative-fill cup game on $n = 2p$ cups. There is an oblivious filling strategy **flatalg** that achieves an R_Δ -flat configuration of cups against a Δ -greedy-like emptier in running time $2(R + \lceil(1 + 1/n)(E + S)\rceil)$. Furthermore, flatalg guarantees that the cup configuration is R -flat on every round.

Proof. If $R \leq R_\Delta$ the algorithm does nothing, since the desired fill-range is already achieved; for the rest of the proof we consider $R > R_\Delta$.

The filler's strategy is to distribute fill equally amongst all cups at every round, placing $p/n = 1/2$ fill in each cup. Let $\ell_t = \min_{c \in S_t} \text{fill}_{S_t}(c)$, $u_t = \max_{c \in S_t} \text{fill}_{S_t}(c)$.

First we show that the fill-range of the cups can only increase if the fill-range is very small.

Claim 7. *If $u_{t+1} - \ell_{t+1} > u_t - \ell_t$ then*

$$u_{t+1} - \ell_{t+1} \leq R.$$

Proof. First we remark that the fill of any cup changes by at most $1/2$ from round to round, and in particular $|u_{t+1} - u_t| \leq 1/2$, $|\ell_{t+1} - \ell_t| \leq 1/2$. In order for the fill-range to increase the emptier must have emptied from some cup with fill in $[\ell_t, \ell_t + 1]$ without emptying from some cup with fill in $[u_t - 1, u_t]$; if the emptier had not emptied from every cup with fill in $[\ell_t, \ell_t + 1]$ then we would have $\ell_{t+1} = \ell_t + 1/2$ so the fill-range cannot have increased, and similarly if the emptier had emptied from every cup with fill in $[u_t - 1, u_t]$ then we would have $u_{t+1} = u_t - 1/2$ so again the fill-range could not have increased. Because the emptier is Δ -greedy-like emptying from a cup with fill at most $\ell_t + 1$ and not emptying from a cup with fill at least $u_t - 1$ implies that $u_t - 1$ and $\ell_t + 1$ differ by at most Δ . Thus,

$$u_{t+1} - \ell_{t+1} \leq u_t + 1/2 - (\ell_t - 1/2) \leq \Delta + 3 \leq R.$$

□

Because by Claim 7 whenever the fill-range of the cups increases it increases to a value at most R , we have by induction that the fill-range of the cups never exceeds R , i.e. the cups are always R -flat.

Let L_t be the set of cups c with $\text{fill}_{S_t}(c) \leq \ell_t + 2 + \Delta$, and let U_t be the set of cups c with $\text{fill}_{S_t}(c) \geq u_t - 2 - \Delta$.

Now we prove a key property of the sets U_t and L_t : if a cup is in U_t or L_t it is also in $U_{t'}, L_{t'}$ for all $t' > t$. This follows immediately from Claim 8.

Claim 8.

$$U_t \subseteq U_{t+1}, \quad L_t \subseteq L_{t+1}.$$

Proof. Consider a cup $c \in U_t$.

If c is not emptied from, i.e. $\text{fill}(c)$ has increased by $1/2$ from the previous round, then clearly $c \in U_{t+1}$, because backlog has increased by at most $1/2$, so $\text{fill}(c)$ must still be within $2 + \Delta$ of the backlog on round $t + 1$.

On the other hand, if c is emptied from, i.e. $\text{fill}(c)$ has decreased by $1/2$, we consider two cases.

Case 1: If $\text{fill}_{S_t}(c) \geq u_t - \Delta - 1$, then $\text{fill}_{S_t}(c)$ is at least 1 above the bottom of the interval defining which cups belong to U_t . The backlog increases by at most $1/2$ and the fill of c decreases by $1/2$, so $\text{fill}_{S_{t+1}}(c)$ is at least $1 - 1/2 - 1/2 = 0$ above the bottom of the interval, i.e. still in the interval.

Case 2: On the other hand, if $\text{fill}_{S_t}(c) < u_t - \Delta - 1$, then every cup with fill in $[u_t - 1, u_t]$ must have been emptied from because the emptier is Δ -greedy-like. Therefore the fullest cup on round $t + 1$ is the same as the fullest cup on round t , because every cup with fill in $[u_t - 1, u_t]$ has had its fill decrease by $1/2$, and no cup with fill less than $u_t - 1$ had its fill increase by more than $1/2$. Hence $u_{t+1} = u_t - 1/2$. Because both $\text{fill}(c)$ and the backlog have decreased by $1/2$, the distance between them is still at most $\Delta + 2$, hence $c \in U_{t+1}$.

The argument for why $L_t \subseteq L_{t+1}$ is symmetric. □

Now we show that under certain conditions u_t decreases and ℓ_t increases.

Claim 9. *On any round t where the emptier empties from at least $n/2$ cups, if $|U_t| \leq n/2$ then $u_{t+1} = u_t - 1/2$. On any round t where the emptier empties from at most $n/2$ cups, if $|L_t| \leq n/2$ then $\ell_{t+1} = \ell_t + 1/2$.*

Proof. Consider a round t where the emptier empties from at least $n/2$ cups. If there are at least $n/2$ cups outside of U_t , i.e. cups with fills in $[\ell_t, u_t - 2 - \Delta]$, then all cups with fills in $[u_t - 2, u_t]$ must be emptied from; if one such cup was not emptied from then by the pigeon-hole principle some cup outside of U_t was emptied from, which is impossible as the emptier is Δ -greedy-like. This clearly implies that $u_{t+1} = u_t - 1/2$: no cup with fill less than $u_t - 2$ has gained enough fill to become the fullest cup, and the fullest cup from the previous round has lost $1/2$ unit of fill.

By a symmetric argument $\ell_{t+1} = \ell_t + 1/2$ if the emptier empties at most $n/2$ cups on a round t where $|L_t| \leq n/2$. □

Now we show that eventually $L_t \cap U_t \neq \emptyset$.

Claim 10. *There is a round $t_0 \leq 2(R + \lceil(1 + 1/n)(E + S)\rceil)$ such that $U_t \cap L_t \neq \emptyset$ for all $t \geq t_0$.*

Proof. We call a round where the emptier doesn't use $p = n/2$ resources, i.e. a round where the number of skipped emptyings and the number of extra emptyings are not equal, an **unbalanced round**; we call a round that is not unbalanced a **balanced round**.

Note that there are clearly at most $E + S$ unbalanced rounds. We now associate some unbalanced rounds with balanced rounds; in particular we define what it means for a balanced round to **cancel** an unbalanced round. We define cancellation by a sequential process. For $i = 1, 2, \dots, 2(R + \lceil (1 + 1/n)(E + S) \rceil)$ (iterating in ascending order of i), if round i is unbalanced then we say that the first balanced round $j > i$ that hasn't already been assigned (earlier in the sequential process) to cancel another unbalanced round $i' < i$, if any such round j exists, **cancels** round i . Note that cancellation is a one-to-one relation: each unbalanced round is cancelled by at most one balanced round and each balanced round cancels at most one unbalanced round.

Consider rounds of the form $2(R + \lceil (E + S)/n \rceil) + (E + S) + i$ for $i \in [E + S + 1] - 1$. We claim that there is some such i such that among rounds $[2(R + \lceil (E + S)/n \rceil) + (E + S) + i]$ every unbalanced round has been cancelled, and such that there are $2(R + \lceil (E + S)/n \rceil)$ balanced rounds not cancelling other rounds. Assume for contradiction that such an i does not exist. Note that there are at least $2(R + \lceil (E + S)/n \rceil)$ balanced rounds in the first $2(R + \lceil (E + S)/n \rceil) + (S + E)$ rounds. Thus every balanced round $2R + (E + S) + \lceil (E + S)/n \rceil + i - 1$ for $i \in [E + S + 1]$ is necessarily a cancelling round, or else there would be a round by which there are no uncanceled unbalanced rounds. Hence by round $2(R + \lceil (E + S)/n \rceil) + 2(E + S)$, there must have been $E + S$ cancelled rounds, so on round $2(R + \lceil (E + S)/n \rceil) + 2(E + S)$ all unbalanced rounds are cancelled, which leaves $2(R + \lceil (E + S)/n \rceil)$ balanced rounds that are not cancelling any rounds, as desired.

Let t_e be the first round by which there are $2(R + \lceil (E + S)/n \rceil)$ balanced non-cancelling rounds. Note that the average fill of the cups cannot have decreased by more than E/n from its starting value; similarly the average fill of the cups cannot have increased by more than S/n . Because the cups start R -flat, we have that u_t cannot have decreased by more than $R + E/n$ or else u_t would necessarily be below the average fill, and identically ℓ_t cannot have increased by more than $R + S/n$ or else it would be above the average fill. Now, by Claim 9 we have that eventually $|L_t| > n/2$: If $|L_t| \leq n/2$ were always true, then on every balanced round ℓ_t would have increased by $1/2$, and since ℓ_t increases by at most $1/2$ on unbalanced rounds, this implies that in total ℓ_t would

have increased by at least $(1/2)2(R + \lceil (E + S)/n \rceil)$, which is impossible. By a symmetric argument it is impossible that $|U_t| \leq n/2$ for all rounds.

Since $|U_{t+1}| \geq |U_t|$ and $|L_{t+1}| \geq |L_t|$ by Claim 8, we have that there is some round $t_0 \in [2(R + \lceil (1 + 1/n)(E + S) \rceil)]$ such that for all $t \geq t_0$ we have $|U_t| > n/2$ and $|L_t| > n/2$. But then we have $U_t \cap L_t \neq \emptyset$, as desired. \square

If there exists a cup $c \in L_t \cap U_t$, then

$$\text{fill}(c) \in [u_t - 2 - \Delta, u_t] \cap [\ell_t, \ell_t + 2 + \Delta].$$

Hence we have that

$$\ell_t + 2 + \Delta \geq u_t - 2 - \Delta.$$

Rearranging,

$$u_t - \ell_t \leq 2(2 + \Delta) = R_\Delta.$$

Thus the cup configuration is R_Δ -flat by the end of this flattening process. \square

Next we describe a simple oblivious filling strategy, that we call **randalg**, that will be used as a subroutine in Lemma 3; this strategy is very well-known, and similar versions of it can be found in [1, 2, 3, 4].

Proposition 3. *Consider an R -flat cup configuration in the single-processor ∞ -extra-emptyings ∞ -skip-emptyings negative-fill cup game on n cups with initial average fill μ_0 . Let $k \in [n]$ be a parameter. Let $d = \sum_{i=2}^k 1/i$.*

*There is an oblivious filling strategy **randalg**(k) with running time $k - 1$ that, if we condition on the emptier not performing extra emptyings, achieves fill at least $\mu_0 - R + d$ in a known cup c with probability at least $1/(n - 1)!$. This filling strategy achieves fill at most $\mu_0 + R + d$ in this cup, unconditionally.*

*Furthermore, when applied against a Δ -greedy-like emptier with $R = R_\Delta$, **randalg**(k) guarantees that the cup configuration is $(R + d)$ -flat on every round.*

Proof. First we condition on the emptier does not using extra emptying and show that in this case the filler has probability at least $1/(k - 1)!$ of attaining a cup with fill at least $\mu_0 - R + d$. The filler maintains an **active set**, initialized to being an arbitrary subset of k of the cups. Every round the filler distributes 1 unit of fill equally among all cups in the active set. Next the emptier removes 1 unit of fill from some cup, or skips its emptying. Then the filler removes a random cup from the active set (chosen uniformly at random from the active set). This continues until a single cup c remains in the active set.

We now bound the probability that c has never been emptied from. Assume that on the i -th step of this process, i.e. when the size of the active set is $n-i+1$, no cups in the active set have ever been emptied from; consider the probability that after the filler removes a cup randomly from the active set there are still no cups in the active set that the emptier has emptied from. If the emptier skips its emptying on this round, or empties from a cup not in the active set then it is trivially still true that no cups in the active set have been emptied from. If the cup that the emptier empties from is in the active set then with probability $1/(k-i+1)$ it is evicted from the active set, in which case we still have that no cup in the active set has ever been emptied from. Hence with probability at least $1/(k-1)!$ the final cup in the active set, c , has never been emptied from. In this case, c will have gained fill $d = \sum_{i=2}^k 1/i$ as claimed. Because c started with fill at least $-R + \mu_0$, c now has fill at least $-R + d + \mu_0$.

Now note that regardless of if the emptier uses extra emptyings c has fill at most $\mu_0 + R + d$, as c starts with fill at most R , and c gains at most $1/(k-i+1)$ fill on the i -th round of this process.

Now we analyze this algorithm specifically for a Δ -greedy-like emptier. Consider a round t on which $\min_c \text{fill}_{S_{t+1}}(c) < \min_c \text{fill}_{S_t}(c)$, and where a cup c_0 that has $\text{fill}_{S_{t+1}}(c_0) = \max_c \text{fill}_{S_{t+1}}(c)$ was not emptied from on round t . Because the emptier is Δ -greedy-like this implies that $\text{fill}_{I_t}(c_0) - \min_c \text{fill}_{I_t}(c) \leq \Delta + 1$ and then $\max_c \text{fill}_{S_{t+1}}(c) - \min_c \text{fill}_{S_{t+1}}(c) \leq \Delta + 2$, i.e. the cups are $(\Delta + 2)$ -flat.

Consider some round t_1 on which the cups are not $(\Delta + 2)$ -flat; let t_0 be the last round on which the cups were R -flat (note that if the cups are $(\Delta + 2)$ -flat they are also R -flat as $\Delta + 2 < R$). Consider how the fill-range of the cups changes during the set of rounds t with $t_0 < t \leq t_1$. On any such round t either $\min_c \text{fill}_{S_{t+1}}(c) \geq \min_c \text{fill}_{S_t}(c)$ in which case the fill-range increases by at most $1/(k-t+1)$ where $k-t+1$ is the size of the active set on round t , or all cups on round $t+1$ with fill equal to the backlog were emptied from, meaning that backlog decreased by at least $1 - 1/(k-t+1)$. In either case the fill-range increases by at most $1/(k-t+1)$. Thus in total the fill-range is at most $R + d$. That is, the cups are $(R + d)$ -flat on round t_1 , as desired. \square

We now give a method for transforming a filling strategy for achieving large backlog into a filling strategy for achieving high fill in many cups, or high average fill in a set of cups (which of these we guarantee depends on the original filling strategy). The idea

of repeating an algorithm many times is also used in the proof of the Adaptive Amplification Lemma; the construction is slightly more complicated in the randomized case however, and is much harder to analyze.

Definition 1. Let alg_0 be an oblivious filling strategy, that can get high fill in some cup against greedy-like emptiers with some probability. We construct a new filling strategy $\text{rep}_\delta(\text{alg}_0)$, which we call the δ -**repetition** of alg_0 , as follows:

Let $n_A = \lceil \delta n \rceil$, $n_B = \lfloor (1 - \delta)n \rfloor$. Let $M \gg n$ be large, let $m = \text{poly}(M)$ be a chosen parameter. Initialize A to \emptyset and B to being all of the cups. We call A the **anchor set** and B the **non-anchor set**. The filler always places 1 unit of fill in each anchor cup on each round. The filling strategy consists of n_A **donation-processes**, which are procedures that result in a cup being **donated** from B to A (i.e. removed from B and added to A). At the start of each donation-process the filler chooses a value m_0 uniformly at random from $[m]$. We say that the filler **applies** a filling strategy alg to B if the filler uses alg on B while placing 1 unit of fill in each anchor cup. During the donation-process the filler applies alg_0 to B m_0 times, and flattens B by applying flatalg to B for $\Theta(M)$ rounds before each application of alg_0 . At the end of each donation process the filler takes the cup given by the final application of alg_0 (i.e. the cup that alg_0 guarantees with some probability against a certain class of emptiers to have a certain high fill), and donates this cup to A .

A pseudocode description of this algorithm can be found in Algorithm 1.

Algorithm 1 rep

Input: $\text{alg}_0, \delta, M, m$, set of n cups

Output: Guarantees on the sets A, B (will vary based on alg_0)

$n_A \leftarrow \lceil \delta n \rceil, n_B \leftarrow \lfloor (1 - \delta)n \rfloor$

$A \leftarrow \emptyset, B \leftarrow$ all the cups

for $i \in [n_A]$ **do** \triangleright Donation-processes

$m_0 \leftarrow \text{random}([m])$

for $j \in [m_0]$ **do**

Apply flatalg to B for $\Theta(M)$ rounds

Apply alg_0 to B

Donate the cup given by alg_0 from B to A

We use the idea of repeating an algorithm in two different contexts. First in Proposition 4 we prove a result analogous to that of Proposition 1: in particular, we show that we can achieve constant fill in a known cup by using $\text{rep}_{\Theta(1)}(\text{randalg}(\Theta(1)))$ which

achieves, by a Chernoff bound, $\Theta(n)$ unknown cups with constant fill, and then exploiting the emptier's greedy-like nature to achieve constant fill in a known cup. After doing this, we prove the **Oblivious Amplification Lemma**, a result analogous to the Adaptive Amplification Lemma: in particular, we show how to take an algorithm for achieving some backlog, and then achieve higher backlog by repeating the algorithm many times. Although these results have deterministic analogues, their proofs are different and significantly more complex than the proofs for the deterministic cases.

Before proving Proposition 4 we analyze $\text{rep}_{\Theta(1)}(\text{randalg}(\Theta(1)))$ in Lemma 3.

Lemma 3. *Let $\Delta \leq O(1)$, let $h \leq O(1)$ with $h \geq 16 + 16\Delta$, let $k = \lceil e^{2h+1} \rceil$, let $\delta = \Theta(e^{-2h})$, let n be at least a sufficiently large constant determined by h and Δ . Let $M \gg n$ be very large. Consider an R_Δ -flat cup configuration in the variable-processor cup game on n cups with initial average fill μ_0 .*

When applied to a Δ -greedy-like emptier $\text{rep}_\delta(\text{randalg}(k))$ either achieves mass at least M in the cups, or with probability at least $1 - 2^{-\Omega(n)}$ makes an (unknown) set of $\Theta(n)$ cups in A have fill at least $h + \mu_0$ while also guaranteeing that $\mu(B) \geq -h/2 + \mu_0$, where A, B are the sets defined in Definition 1. This strategy has running time $\text{poly}(M)$.

Proof. We use all definitions given in the Definition 1.

Without loss of generality we assume that the emptier does not neglect the anchor set more than M in a particular donation process; if the emptier chooses to neglect the anchor set this much then the anchor cups will have achieved mass M so Lemma 3 is already fulfilled. Similarly we assume that the emptier does not choose to skip more than M emptyings.

We say that the emptier **neglects** the anchor set on a round if it does not empty from each anchor cup. We say that an application of $\text{randalg}(k)$ to B is **non-emptier-wasted** if the emptier does not neglect the anchor set during any round of the application of $\text{randalg}(k)$. We define $d = \sum_{i=2}^k 1/i = \Theta(h)$. We say that an application of $\text{randalg}(k)$ to D is **lucky** if it achieves backlog at least $\mu(B) - R_\Delta + d$ where $\mu(B)$ is measured at the start of the application of $\text{randalg}(k)$; note that by Proposition 3 if we condition on an application of $\text{randalg}(k)$ where B started R_Δ -flat being non-emptier-wasted then the application has at least a $1/k!$ chance of being lucky.

Now we prove several important bounds on fills of cups in A and B .

Claim 11. *All applications of flatalg make B be R_Δ -flat and B is always $(R_\Delta + d)$ -flat.*

Proof. Given that the application of flatalg immediately prior to an application of $\text{randalg}(k)$ made B be R_Δ -flat, by Proposition 3 we have that B will stay $(R_\Delta + d)$ -flat during the application of $\text{randalg}(k)$. Given that the application of $\text{randalg}(k)$ immediately prior to an application of flatalg resulted in B being $(R_\Delta + d)$ -flat, we have that B remains $(R_\Delta + d)$ -flat throughout the duration of the application of flatalg by Lemma 2. Given that B is $(R_\Delta + d)$ -flat before a donation occurs B is clearly still $(R_\Delta + d)$ -flat after the donation, because the only change to B during a donation is that a cup is removed from B which cannot increase the fill-range of B . Note that B started R_Δ -flat at the beginning of the first donation-process. Note that if an application of flatalg begins with B being $(R_\Delta + d)$ -flat, then by considering the flattening to happen in the $(|B|/2)$ -processor M -extra-emptyings M -skip-emptyings cup game we ensure that it makes B be R_Δ -flat. Hence we have by induction that B has always been $(R_\Delta + d)$ -flat and that all flattening processes have made B be R_Δ -flat. \square

Now we aim to show that $\mu(B)$ is never very low, which we need in order to establish that every non-emptier-wasted lucky application of $\text{randalg}(k)$ gets a cup with high fill. Interestingly in order to lower bound $\mu(B)$ we first must upper bound $\mu(B)$, which by greediness and flatness of B gives an upper bound on $\mu(A)$ which we use to get a lower bound on $\mu(B)$.

Claim 12. *We have always had*

$$\mu(B) \leq 2 + \mu(AB).$$

Proof. There are two ways that $\mu(B) - \mu(AB)$ can increase:

Case 1: The emptier could empty from 0 cups in B while emptying from every cup in A .

Case 2: The filler could evict a cup with fill lower than $\mu(B)$ from B at the end of a donation-process.

Note that cases are exhaustive, in particular note that if the emptier skips more than 1 emptying then $\mu(B) - \mu(AB)$ must decrease because $|B| \approx |AB|$, in particular by our choice of $\delta = \Theta(e^{-2h})$, as opposed to in Case 1 where $\mu(B) - \mu(AB)$ increases.

In Case 1, because the emptier is Δ -greedy-like,

$$\min_{a \in A} \text{fill}(a) > \max_{b \in B} \text{fill}(b) - \Delta.$$

Thus $\mu(B) \leq \mu(A) + \Delta$. As $|B| \gg |A|$, in particular by our choice of $\delta = \Theta(e^{-2h})$, this can be loosened to $\mu(B) \leq 1 + \mu(AB)$.

Consider the final round on which B is skipped while A is not skipped (or consider the first round if there is no such round).

From this round onwards the only increase to $\mu(B) - \mu(AB)$ is due to B evicting cups with fill well below $\mu(B)$. We can upper bound the increase of $\mu(B) - \mu(AB)$ by the increase of $\mu(B)$ as $\mu(AB)$ is strictly increasing.

The cup that B evicts at the end of a donation-process has fill at least $\mu(B) - R_\Delta - (k-1)$, as the running time of $\text{randalg}(k)$ is $k-1$, and because B starts R_Δ -flat by Claim 11. Evicting a cup with fill $\mu(B) - R_\Delta - (k-1)$ from B changes $\mu(B)$ by $(R_\Delta + k - 1)/(|B| - 1)$ where $|B|$ is the size of B before the cup is evicted from B . Even if this happens on each of the n_A donation-processes $\mu(B)$ cannot rise higher than $n_A(R_\Delta + k - 1)/(n - n_A)$ which by design in choosing $n_B \gg n_A$, as was done in choosing $\delta = \Theta(e^{-2h})$, is at most 1.

Thus $\mu(B) \leq 2 + \mu(AB)$ is always true. \square

The upper bound on $\mu(B)$ along with the guarantee that B is flat allows us to bound the highest that a cup in A could rise by greediness, which in turn upper bounds $\mu(A)$ which in turn lower bounds $\mu(B)$. In particular we have

Claim 13. *We always have*

$$\mu(B) \geq -h/2 + \mu_0.$$

Proof. By Claim 12 and Claim 11 we have that no cup in B ever has fill greater than $u_B = \mu(AB) + 2 + R_\Delta + d$. Let $u_A = u_B + \Delta + 1$. We claim that the backlog in A never exceeds u_A . Note that $\mu(AB), u_A, u_B$ are implicitly functions of the round; $\mu(AB)$ can increase from μ_0 if the emptier skips emptyings.

Consider how high the fill of a cup $c \in A$ could be. If c came from B then when it is donated to A its fill is at most $u_B < u_A$. Otherwise, c started with fill at most $R_\Delta < u_A$. Now consider how much the fill of c could increase while being in A . Because the emptier is Δ -greedy-like, if a cup $c \in A$ has fill more than Δ higher than the backlog in B then c must be emptied from, so any cup with fill at least $u_B + \Delta = u_A - 1$ must be emptied from, and hence u_A upper bounds the backlog in A .

Of course an upper bound on backlog in A also serves as an upper bound on the average fill of A as well, i.e. $\mu(A) \leq u_A$. Rearranging the expression

$$|B|\mu(B) + |A|\mu(A) = |AB|\mu(AB)$$

we have

$$\begin{aligned} \mu(B) &= -\frac{|A|}{|B|}\mu(A) + \frac{|AB|}{|B|}\mu(AB) \\ &\geq -(\mu(AB) + 3 + R_\Delta + d + \Delta)\frac{|A|}{|B|} + \frac{|AB|}{|B|}\mu(AB) \\ &= -(3 + R_\Delta + d + \Delta)\frac{|A|}{|B|} + \mu(AB) \\ &\geq -h/2 + \mu(AB) \end{aligned}$$

where the final inequality follows because $\mu(AB) \geq 0$, and $|B| \gg |A|$, in particular by our choice of $\delta = \Theta(e^{-2h})$. Of course $\mu(AB) \geq \mu_0$ so we have

$$\mu(B) \geq -h/2 + \mu_0. \quad \square$$

Now we show that at least a constant fraction of the donation-processes succeed with exponentially good probability.

Claim 14. *There exists choice of $m = \Theta(M)$ such that with probability at least $1 - 2^{-\Omega(n)}$, the filler achieves fill at least $h + \mu_0$ in $\Theta(n)$ of the cups in A .*

Proof. If the emptier was not allowed to neglect the anchor set ever then the claim would clearly be true as each application of $\text{randalg}(k)$ would unconditionally succeed with constant probability, so a Chernoff bound would give that $\Theta(n)$ of the donation-processes donate a cup with fill at least $\mu(B) - R_\Delta + d \geq h + \mu_0$, where the inequality follows from Claim 13 which asserts that $\mu(B) \geq -h/2 + \mu_0$, and from the facts $d \geq 2h$ and $h \geq 16(1 + \Delta)$. However, the emptier is allowed to neglect the anchor set, and worse, the emptier can choose to neglect the anchor set conditional on the filler's progress during $\text{randalg}(k)$. However, by applying $\text{randalg}(k)$ a random number of times, chosen from $[m]$ (where $m = \Theta(M)$ which is quite large), we guarantee that with exponentially good probability the filler succeeds many times, in particular $\Theta(M)$ times. But since the emptier cannot neglect the anchor set more than M times, by appropriately large choice of m we can make it so that the filler succeeds at least $2M$ times with exponentially good probability. Then the emptier would have at best a $1/2$ chance of preventing the donation-process from giving away a cup with fill $h + \mu_0$ whenever one such cup is achieved. We now formalize this reasoning.

We can lower bound the probability of getting $\Theta(n)$ cups with fills all at least $h + \mu_0$ by considering an augmented emptier that is allowed to interfere with

M applications of $\text{randalg}(k)$ per donation-process that only interferes with applications of $\text{randalg}(k)$ that would otherwise donate a cup with fill at least $h + \mu_0$ into A . The optimal strategy for such an emptier, given our filler's strategy of randomly choosing how many times to apply $\text{randalg}(k)$ before donating a cup, is to simply interfere with the first M applications of $\text{randalg}(k)$ that would have achieved a cup with fill $h + \mu_0$ without interference. Let $m = 4Mk! = \Theta(M)$. Recall that conditional on the emptier not interfering, each of these applications of $\text{randalg}(k)$ has at least a $1/k!$ chance of getting a cup with fill h . Hence, by a Chernoff bound with exponentially good probability at least $2M$ of m applications of $\text{randalg}(k)$ have the potential to donate a cup with fill $h + \mu_0$ to A , if the emptier does not interfere. The filler chooses an application uniformly at random from $[m]$ on which to donate a cup. With probability at least $1/k!$ this is on an application where the filler could get a cup with fill $h + \mu_0$ in A if the emptier does not interfere, and with probability at least $1/2$ the emptier does not interfere on this application of $\text{randalg}(k)$, because the emptier can interfere on at most M of the applications of $\text{randalg}(k)$.

Against this augmented emptier whether or not donation-processes achieve a cup with fill $h + \mu_0$ in A are independent events. As each happens with at least constant probability, by a Chernoff bound there is exponentially high probability that at least a constant fraction of them succeed.

Note that we used the Chernoff bound $\Theta(n)$; by a union bound there is exponentially good probability that all of the desired events occur. \square

We now analyze the running time of the filling strategy. There are n_A donation-processes. Each donation-process consists of $\Theta(M)$ applications of $\text{randalg}(k)$, which each take constant time, and $\Theta(M)$ applications of flatalg , which each take $\text{poly}(M)$ time. Thus overall the algorithm takes $\text{poly}(M)$ time, as desired. \square

Now, using Lemma 3 we show in Proposition 4 that an oblivious filler can achieve constant backlog.

Proposition 4. *Let $H \leq O(1)$, let $\Delta \leq O(1)$, let n be at least a sufficiently large constant determined by H and Δ . Let $M \gg n$ be very large. Consider an R_Δ -flat cup configuration in the variable-processor cup game on n cups with average fill μ_0 . There is an oblivious filling strategy that either achieves mass M among the cups, or achieves fill at least $\mu_0 + H$ in*

a chosen cup in running time $\text{poly}(M)$ against a Δ -greedy-like emptier with probability at least $1 - 2^{-\Omega(n)}$.

Proof. The filler starts by using $\text{rep}_\delta(\text{randalg}(k))$ with parameter settings as in Lemma 3 where $h = H \cdot 16(1 + \Delta)$, i.e. $k = \lceil e^{2h+1} \rceil$, $\delta = \Theta(e^{-2h})$. Let the number of cups which must now exist by Lemma 3 with fill at least $h + \mu_0$ be of size $nc = \Theta(n)$.

The filler sets $p = 1$, i.e. uses a single processor. Now the filler exploits the emptier's greedy-like nature to get fill H in a chosen cup c_0 . Specifically, for $(5/8)h$ rounds the filler places 1 unit of fill into c_0 . Because the emptier is Δ -greedy-like it must empty from the nc cups in A with fill at least $h + \mu_0$ until c_0 has large fill. Over $(5/8)h$ rounds the cups in A cannot have their fill decrease below $(3/8)h \geq h/8 + \Delta + \mu_0$. Hence, any cups with fills less than $h/8 + \mu_0$ must not be emptied from during these rounds. The fill of c_0 started as at least $-h/2 + \mu_0$ as $\mu(B) \geq -h/2 + \mu_0$. After $(5/8)h$ rounds c_0 has fill at least $h/8 + \mu_0$, because the emptier cannot have emptied c_0 until it attained fill $h/8 + \mu_0$, and if c_0 is never emptied from then it achieves fill $h/8 + \mu_0$. Thus the filling strategy achieves backlog $h/8 + \mu_0 \geq H + \mu_0$ in c_0 , a known cup, as desired. \square

Next we prove the **Oblivious Amplification Lemma**.

Lemma 4 (Oblivious Amplification Lemma). *Let $\delta \in (0, 1/2)$ be a constant parameter. Let $\Delta \leq O(1)$. Let M be very large. Consider a cup configuration in the variable-processor cup game on $n \ll M, n > \Omega(1/\delta^2)$ cups with average fill μ_0 that is R_Δ -flat. Let $\text{alg}(f)$ be an oblivious filling strategy that either achieves mass M or achieves backlog $\mu_0 + f(n)$ on such cups with probability at least $1 - 2^{-\Omega(n)}$ in running time $T(n)$ against a Δ -greedy-like emptier.*

Consider a cup configuration in the variable-processor cup game on $n \ll M, n > \Omega(1/\delta^2)$ cups with average fill μ_0 that is R_Δ -flat. There exists an oblivious filling strategy $\text{alg}(f')$ that either achieves mass M or achieves backlog $f'(n)$ satisfying

$$f'(n) \geq (1 - \delta)^2 f(\lfloor (1 - \delta)n \rfloor) + f(\lceil \delta n \rceil) + \mu_0$$

and $f'(n) \geq f(n)$, with probability at least $1 - 2^{-\Omega(n)} - 1/\text{poly}(M)$ in running time

$$T'(n) \leq M \cdot n \cdot T(\lfloor (1 - \delta)n \rfloor) + T(\lceil \delta n \rceil)$$

against a Δ -greedy-like emptier.

Proof. We use the notation from Lemma 3, and from Definition 1. To summarize, we define $n_A =$

$\lceil \delta n \rceil, n_B = \lfloor (1 - \delta)n \rfloor$, we refer to A as the anchor set and B as the non-anchor set, we say that the filler applies $\text{alg}(f)$ to B if it uses $\text{alg}(f)$ on B while placing 1 fill into each cup in A , and we say that A is neglected during an application of $\text{alg}(f)$ to B if there is some round during the application where the emptier does not empty from all anchor cups.

The filler defaults to using $\text{alg}(f)$ on all the cups if

$$f(n) \geq (1 - \delta)^2 f(n_B) + f(n_A).$$

In this case our strategy trivially has the desired guarantees. In the rest of the proof we consider the case where we cannot simply fall back on $\text{alg}(f)$ to achieve the desired backlog.

The filler's strategy is roughly as follows:

Step 1: Use $\text{rep}_\delta(\text{alg}(f))$ on all the cups; this will get A to have high average fill.

Step 2: Flatten A using flatalg , and then use $\text{alg}(f)$ on A .

Now we analyze Step 1, and show that by appropriately choosing parameters it can be made to succeed.

Note that, exactly as in the proof of Lemma 3, the emptier cannot neglect the anchor set more than M times per donation-process, and the emptier cannot skip more than M emptyings, without causing the mass of the cups to be at least M ; we assume the emptier chooses not to do this.

We choose $m \geq M^3, m = \text{poly}(M)$ —recall that $[m]$ is the set from which we choose how many times to apply $\text{alg}(f)$ in a donation-process. Using the ideas from the analysis in Claim 14, we see that by a Chernoff bound with exponentially good probability in n many more than $m/2$ of the applications would succeed without if the emptier did not interfere with them. Conditional on the final application being a round that would be successful without the emptier interfering, the emptier has at best a $M/(m/2) = 1/\text{poly}(M)$ chance of interfering on the correct round. Taking a union bound, we can say that with probability at least $1 - 1/\text{poly}(M) - 2^{-\Omega(n)}$ all applications of $\text{alg}(f)$ are not non-emptier-wasted and successfully achieve a cup with fill at least $\mu_{t_0}(B) + f(n_B)$ where $\mu_{t_0}(B)$ refers to the average fill of B measured at the start of the application of $\text{alg}(f)$.

Let skips_t denote the number of times that the emptier has skipped the anchor set by step t . Consider how $\mu(B) - \text{skips}/n_B$ changes over the course of the donation processes. As noted above at the end of each donation-process $\mu(B)$ decreases due to B donating a cup with fill at least $\mu(B) + f(n_B)$. In particular, if t_0 denotes the time right before a cup is donated on the i -th donation-process and t_1 denotes the time right after a cup is donated, then $\mu_{t_1}(B) = \mu_{t_0}(B) - f(n_B)/(n - i)$. Now we claim that

$\mu(B) - \text{skips}/n_B$ is monotonically decreasing. Clearly each donation decreases this quantity. If the anchor set is neglected then $\mu(B)$ decreases. If a skip occurs, then skips/n_B increases by more than $\mu(B)$ can possibly decrease. Let t_* be the time at the end of all the donation-processes. We have that

$$\mu_{t_*}(B) - \frac{\text{skips}_{t_*}}{n_B} \leq \mu_0 - \sum_{i=1}^{n_A} \frac{f(n_B)}{n - i}. \quad (12)$$

By conservation of mass we have

$$n_A \cdot \mu_{t_*}(A) + n_B \cdot \mu_{t_*}(B) = \mu_0 + \text{skips}_{t_*}. \quad (13)$$

We can use Inequalities (13) and (12) to get a lower bound on $\mu_{t_*}(A)$ as follows:

$$\mu_{t_*}(A) = \mu_0 + \frac{n_B}{n_A} \left(\mu_0 + \frac{\text{skips}_{t_*}}{n_B} - \mu_{t_*}(B) \right). \quad (14)$$

Now we obtain a simpler form of Inequality (12). Let H_n denote the n -th harmonic number. We desire a simpler lower bound for

$$\sum_{i=1}^{n_A} \frac{1}{n - i} = H_{n-1} - H_{n_B-1}.$$

We use the well known fact that

$$\frac{1}{2(n+1)} < H_n - \ln n - \gamma < \frac{1}{2n} \quad (15)$$

where $\gamma = \Theta(1)$ denotes the Euler-Mascheroni constant. Of course $H_{n-1} - H_{n_B-1} \geq H_n - H_{n_B}$. Now using Inequality (15) we have

$$\begin{aligned} H_n - H_{n_B} &> \left(\ln n + \gamma + \frac{1}{2(n+1)} \right) - \left(\ln n_B + \gamma + \frac{1}{2n_B} \right) \\ &> \ln \frac{1}{1 - \delta} + \frac{1}{2} \left(\frac{n_B - n - 1}{(n+1)n_B} \right) \\ &> \delta - \Theta \left(\frac{\delta}{(1 - \delta)n} \right). \end{aligned}$$

Now using this lower bound on $H_n - H_{n_B}$ in Inequality (14) we have:

$$\begin{aligned} \mu_{t_*}(A) &> \mu_0 + \frac{n_B}{n_A} \left(\delta - \Theta \left(\frac{\delta}{(1 - \delta)n} \right) \right) f(n_B) \\ &= \mu_0 + \frac{\lfloor (1 - \delta)n \rfloor}{\lceil \delta n \rceil} \left(\delta - \Theta \left(\frac{\delta}{(1 - \delta)n} \right) \right) f(n_B) \\ &> \mu_0 + \left(\frac{1 - \delta}{\delta} - \frac{1}{\delta^2 n} \right) \left(\delta - \Theta \left(\frac{\delta}{(1 - \delta)n} \right) \right) f(n_B) \\ &> \mu_0 + ((1 - \delta) - \Theta(1/(\delta n))) f(n_B). \end{aligned}$$

Thus, by choosing $n > \Omega(1/\delta^2)$ we have

$$\mu_{t_*}(A) > \mu_0 + (1 - \delta)^2 f(n_B).$$

We have shown that in Step 1 the filler achieves average fill $\mu_0 + (1 - \delta)f(n_B)$ in A (with good probability). Now the filler flattens A and uses $\text{alg}(f)$ on A . It is clear that this is possible, and succeeds with probability at least $1 - 2^{-\Omega(n)}$. This gets a cup with fill

$$\mu_0 + (1 - \delta)^2 f(n_B) + f(n_A)$$

in A , as desired.

Taking a union bound over the probabilities of Step 1 and Step 2 succeeding gives the desired probability.

The running time of Step 1 is clearly $M \cdot n \cdot T(\lfloor (1 - \delta)n \rfloor)$ and the running time of Step 2 is clearly $T(\lceil \delta n \rceil)$ summing these yields the desired upper bound on running time. \square

Finally we prove that an oblivious filler can achieve backlog $n^{1-\epsilon}$, just like an adaptive filler despite the oblivious filler's disadvantage. The proof is very similar to the proof of Theorem 1, but more complicated because in the oblivious case we must guarantee that the result holds with good probability, and also more complicated because the Oblivious Amplification Lemma is more complicated than the Adaptive Amplification Lemma.

Theorem 3. *There is an oblivious filling strategy for the variable-processor cup game on n cups that achieves backlog at least $\Omega(n^{1-\epsilon})$ for any constant $\epsilon > 0$ in running time $2^{O(\log^2 n)}$ with probability at least $1 - 2^{-\Omega(n)}$ against any Δ -greedy-like emptier for $\Delta \leq O(1)$.*

Proof. We aim to achieve backlog $(n/n_b)^{1-\epsilon} - 1$ for some constant n_b on n cups. Let δ be a constant, chosen as a function of ϵ .

By Proposition 4 there is an oblivious filling strategy that achieves backlog $\Omega(1)$ on n cups with exponentially good probability in n ; we call this algorithm $\text{alg } f_0$. However, unlike in the proof of Theorem 1, we obviously cannot use the base case with a constant number of cups: doing so would completely destroy our probability of success. Because the running time of our algorithm will be $2^{\text{polylog}(n)}$, we will be required to take a union bound over $2^{\text{polylog}(n)}$ events. By making the size of our base case $n_b = \text{polylog}(n)$ we get that the probability of the algorithm failing in the base case is at most $2^{-\text{polylog}(n)}$. Then, taking a union bound over $2^{\text{polylog}(n)}$ events can give us the desired probability of success. By Proposition 4 $\text{alg } f_0$ achieves backlog $f_0(k) \geq H \geq \Omega(1)$ for all $k \geq n_b$, for some constant $H \geq \Omega(1)$ to be determined (H is a function of δ).

Then we construct f_{i+1} as the amplification of f_i using Lemma 4.

Define a sequence g_i as

$$g_i = \begin{cases} n_b \lceil 16/\delta \rceil, & i = 0 \\ \lfloor g_{i-1}/(1 - \delta) \rfloor, & i \geq 1 \end{cases}.$$

We claim the following regarding our construction:

Claim 15.

$$f_i(k) \geq (k/n_b)^{1-\epsilon} - 1 \text{ for all } k \leq g_i. \quad (16)$$

Proof. We prove Claim 15 by induction on i .

When $i = 0$, the base case of our induction, (16) is trivially true as $(k/n_b)^{1-\epsilon} - 1 \leq H$ by definition of H for $k \leq g_0$.

Assume (16) for f_i , consider f_{i+1} .

Note that, by design of g_i , if $k \leq g_{i+1}$ then $\lfloor k \cdot (1 - \delta) \rfloor \leq g_i$. Consider any $k \in [g_{i+1}]$.

First we deal with the trivial case where $k \leq g_0$. In this case

$$f_{i+1}(k) \geq f_i(k) \geq \dots \geq f_0(k) \geq (k/n_b)^{1-\epsilon} - 1.$$

Now we consider $k \geq g_0$. Note that in this case $\lfloor (1 - \delta)k \rfloor \geq n_b$. Since f_{i+1} is the amplification of f_i , and k is sufficiently large, we have by Lemma 4 that

$$f_{i+1}(k) \geq (1 - \delta)^2 f_i(\lfloor (1 - \delta)k \rfloor) + f_i(\lceil \delta k \rceil).$$

By our inductive hypothesis, which applies as $\lceil \delta k \rceil \leq g_i$, $\lfloor k \cdot (1 - \delta) \rfloor \leq g_i$, we have

$$f_{i+1}(k) \geq (1 - \delta)^2 (\lfloor (1 - \delta)k/n_b \rfloor^{1-\epsilon} - 1) + \lceil \delta k/n_b \rceil^{1-\epsilon} - 1.$$

Dropping the floor and ceiling, incurring a -1 for dropping the floor, we have

$$f_{i+1}(k) \geq (1 - \delta)^2 ((1 - \delta)k/n_b - 1)^{1-\epsilon} - 1 + (\delta k/n_b)^{1-\epsilon} - 1.$$

Because $(x - 1)^{1-\epsilon} \geq x^{1-\epsilon} - 1$, due to the fact that $x \mapsto x^{1-\epsilon}$ is a sub-linear sub-additive function, we have

$$f_{i+1}(k) \geq (1 - \delta)^2 ((1 - \delta)k/n_b)^{1-\epsilon} - 2 + (\delta k/n_b)^{1-\epsilon} - 1.$$

Moving the $(k/n_b)^{1-\epsilon}$ to the front we have

$$f_{i+1}(k) \geq (k/n_b)^{1-\epsilon} \cdot \left((1 - \delta)^{3-\epsilon} + \delta^{1-\epsilon} - \frac{2(1 - \delta)^2}{(k/n_b)^{1-\epsilon}} \right) - 1.$$

Because $(1 - \delta)^{3-\epsilon} \geq 1 - (3 - \epsilon)\delta$, a fact called Bernoulli's Identity, we have

$$f_{i+1}(k) \geq (k/n_b)^{1-\epsilon} \cdot \left(1 - (3 - \epsilon)\delta + \delta^{1-\epsilon} - \frac{2(1 - \delta)^2}{(k/n_b)^{1-\epsilon}} \right) - 1.$$

Of course $-2(1 - \delta)^2 > -2$, so

$$f_{i+1}(k) \geq (k/n_b)^{1-\epsilon} \cdot (1 - (3 - \epsilon)\delta + \delta^{1-\epsilon} - 2/(k/n_b)^{1-\epsilon}) - 1.$$

Because

$$\frac{-2}{(k/n_b)^{1-\varepsilon}} \geq \frac{-2}{(g_0/n_b)^{1-\varepsilon}} \geq -2(\delta/16)^{1-\varepsilon} \geq -\delta^{1-\varepsilon}/2,$$

which follows from our choice of $g_0 = \lceil 8/\delta \rceil n_b$ and the restriction $\varepsilon < 1/2$, we have

$$f_{i+1}(k) \geq (k/n_b)^{1-\varepsilon} \cdot (1 - (3 - \varepsilon)\delta + \delta^{1-\varepsilon} - \delta^{1-\varepsilon}/2) - 1.$$

Finally, combining terms we have

$$f_{i+1}(k) \geq (k/n_b)^{1-\varepsilon} \cdot (1 - (3 - \varepsilon)\delta + \delta^{1-\varepsilon}/2) - 1.$$

Because $\delta^{1-\varepsilon}$ dominates δ for sufficiently small δ , there is a choice of $\delta = \Theta(1)$ such that

$$1 - (3 - \varepsilon)\delta + \delta^{1-\varepsilon}/2 \geq 1.$$

Taking δ to be this small we have,

$$f_{i+1}(k) \geq (k/n_b)^{1-\varepsilon} - 1,$$

completing the proof. \square

The sequence g_i is n_b times the sequence g_i from the proof of Theorem 1; we thus have that $g_{i_*} \geq n$ for some $i_* \leq O(\log n)$. Hence $\text{alg } f_{i_*}$ achieves backlog

$$f_{i_*}(n) \geq (n/n_b)^{1-\varepsilon} - 1.$$

As $n_b \leq \text{polylog}(n)$ we have

$$f_{i_*}(n) \geq \Omega(n^{1-\varepsilon}),$$

as desired.

Let the running time of $f_i(n)$ be $T_i(n)$. From the Amplification Lemma we have following recurrence bounding $T_i(n)$:

$$\begin{aligned} T_i(n) &\leq 6n^{\eta+1}\delta \cdot T_{i-1}(\lfloor (1-\delta)n \rfloor) + T_{i-1}(\lceil \delta n \rceil) \\ &\leq 7n^{\eta+1}T_{i-1}(\lfloor (1-\delta)n \rfloor). \end{aligned}$$

It follows that $\text{alg } f_{i_*}$, recalling that $i_* \leq O(\log n)$, has running time

$$T_{i_*}(n) \leq (7n^{\eta+1})^{O(\log n)} \leq 2^{O(\log^2 n)}$$

as desired.

As noted, because the running time is $2^{\text{polylog}(n)}$ and the base case size is $n_b \geq \text{polylog}(n)$, a union bound guarantees the probability of success is at least $1 - 2^{-\text{polylog}(n)}$. \square

6 Conclusion

Many important open questions remain open. Can our oblivious cup game results be improved, e.g. by expanding them to apply to a broader class of emptiers? Can the classic oblivious multi-processor cup-game be tightly analyzed?

Maybe more importantly is the implications for analysis of parallel algorithms in general. Are there other problems where letting the number of processors change drastically effects the outcome? This is important to look into.

References

- [1] Michael A Bender, Martín Farach-Colton, and William Kuszmaul. Achieving optimal backlog in multi-processor cup games. In *Proceedings of the 51st Annual ACM SIGACT Symposium on Theory of Computing*, pages 1148–1157, 2019.
- [2] Michael A Bender, Sándor P Fekete, Alexander Kröller, Vincenzo Liberatore, Joseph SB Mitchell, Valentin Polishchuk, and Jukka Suomela. The minimum backlog problem. *Theoretical Computer Science*, 605:51–61, 2015.
- [3] Paul Dietz and Rajeev Raman. Persistence, amortization and randomization. 1991.
- [4] William Kuszmaul. Achieving optimal backlog in the vanilla multi-processor cup game. *SIAM*, 2020.